

# APPLICATIONS



OF DATA SCIENCE

# Intro to Web Scraping

Applications of Data Science - Class Bonus

Giora Simchoni

[gsimchoni@gmail.com](mailto:gsimchoni@gmail.com) and add #dsapps in subject

Stat. and OR Department, TAU

2023-03-26

APPLICATIONS



OF DATA SCIENCE

# The Three Rules of Web Scraping

APPLICATIONS



OF DATA SCIENCE

# Rule 1: Do you *really* need web scraping?

There are data APIs for just about anything, you know...

The image is a collage of nine logos, each representing a different data API or service:

- CHUCKNORRIS.ID**: A logo featuring a cartoon character of Chuck Norris wearing a sombrero and holding a gun, with the text "CHUCKNORRIS.ID" below it.
- PokeAPI**: A logo for "The RESTful Pokémon API" with the text "Serving over 60,000,000 API calls each month!" below it.
- Harvard Art Museums**: Logos for the Fogg Museum, Busch-Reisinger Museum, and Arthur M. Sackler Museum.
- FBI MOST WANTED**: The official logo of the FBI's Most Wanted list.
- FLUTRACK**: A logo with the word "FLUTRACK" in red and a small circular icon.
- ALPHA VANTAGE**: A logo with a dark background featuring a circuit board pattern.
- Spotify**: The Spotify logo with its green circle and white play button icon.
- RainViewer**: A logo with a blue square containing a white cloud-like icon.
- Clinical Table Search Service**: A logo from the U.S. National Library of Medicine, LHN CBC, featuring the NIH logo and the text "A web API service for use with autocomplete-lhc & LHC-Forms".

## R API Packages

Many of them already accessible with a R/Python package...

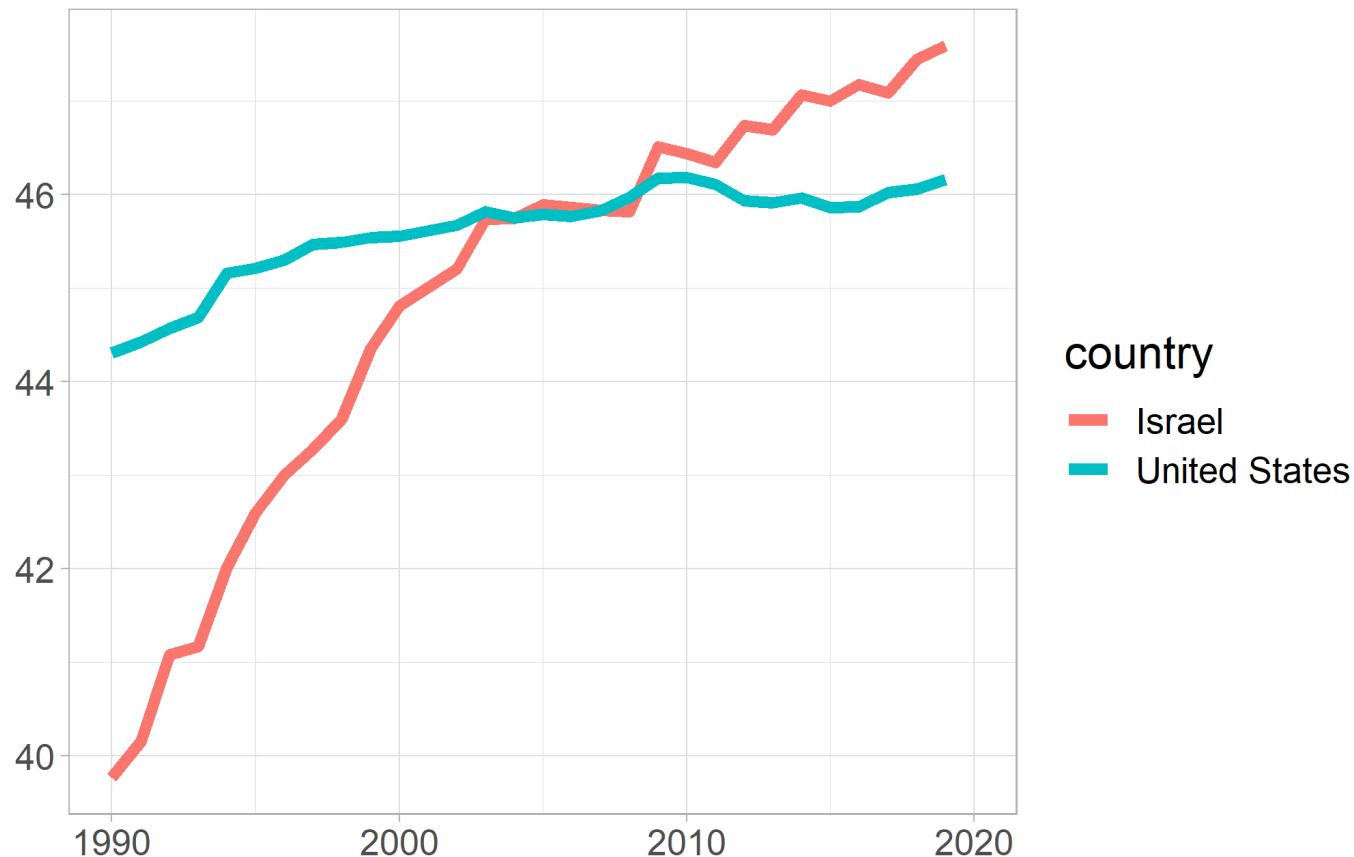
```
library(wbstats)

female_labor <- wb_data(
  indicator = c("women_lab_share" = "SL.TLF.TOTL.FE.ZS"),
  start_date = 1990,
  end_date = 2020
)

female_labor %>%
  filter(country %in% c("Israel", "United States")) %>%
  ggplot(aes(date, women_lab_share, color = country)) +
  geom_line(lwd = 2) +
  labs(title = "Share of women in labor force") +
  theme_light() +
  theme(text = element_text(size=16))
```

From: <https://cfss.uchicago.edu/notes/application-program-interface/>

# % of women in labor force



country  
— Israel  
— United States

## The datapasta package

My gift to you.



# Rule 2: Learn some HTML first!

HTML is a set (or tree) of *elements*, marked by *HTML tags*:

```
<!DOCTYPE html>
<html>
<head>
<title>My Awesome Webpage</title>
</head>
<body>

<h1>My Superb Heading</h1>
<p>I am going to be a web designer.</p>

</body>
</html>
```

## My Superb Heading

I am going to be a web designer.

- First children in the tree: header and body
- View any page's HTML (on chrome) with right-click and "View page source" (or Ctrl + U)

## Useful elements and attributes to know

- `<p> for paragraph </p>`
- `<h1> for headings </h1>`
- `<br>, <hr> for breaks`
- `<a href = "http://www.google.com> for links </a>`
- `<b><i> For bold, italic etc. </i></b>`
- ``
- `<p style="color:DodgerBlue;"> for font color </p>`

# HTML Tables

A big thing when it comes to data as you can imagine...

```
<!DOCTYPE html>
<html>
<body>

<h2>Basic HTML Table</h2>

<table style="width:100%">
  <tr>
    <th>Firstname</th>
    <th>Lastname</th>
    <th>Age</th>
  </tr>
  <tr>
    <td>Britney</td>
    <td>Spears</td>
    <td>39</td>
  </tr>
  <tr>
    <td>Christina</td>
    <td>Aguillera</td>
    <td>40</td>
  </tr>
  <tr>
    <td>Beyonce</td>
    <td>Knowles</td>
    <td>39</td>
  </tr>
</table>

</body>
</html>
```

**Basic HTML Table**

Firstname	Lastname	Age
Britney	Spears	39
Christina	Aguillera	40
Beyonce	Knowles	39

# HTML Classes

A class attribute is defined in a style sheet, lets you repeat a style.

```
<!DOCTYPE html>
<html>
<head>
<style>
.city {
    background-color: tomato;
    color: white;
    border: 2px solid black;
    margin: 20px;
    padding: 20px;
}
</style>
</head>
<body>

<div class="city">
<h2>London</h2>
<p>London is the capital of England.</p>
</div>

<div class="city">
<h2>Paris</h2>
<p>Paris is the capital of France.</p>
</div>

<div class="city">
<h2>Tokyo</h2>
<p>Tokyo is the capital of Japan.</p>
</div>

</body>
</html>
```

London

London is the capital of England.

Paris

Paris is the capital of France.

Tokyo

Tokyo is the capital of Japan.

# Rule 3: Be polite!

With great power comes great responsibility.

See e.g. the [polite](#) package.



# rvest

APPLICATIONS



OF DATA SCIENCE

## **read\_html()**

You're now a NLP expert, and you've just developed a SOTA Q&A model. How would you demonstrate your model's performance?

How about [triviaquestionsnow.com?](https://www.triviaquestionsnow.com/)

Let's scrape a few Q&As. Politely.

```
library(rvest)  
  
url <- "https://www.triviaquestionsnow.com/for/sports-trivia"  
  
html_obj <- read_html(url)
```

`read_html()` is usually where you'd start. What did you get?

```
class(html_obj)  
  
## [1] "xml_document" "xml_node"
```

# View page source

With time, you'll become friendly with this weird object. Right now it is better than...

```
1 <!DOCTYPE html>
2 <html lang="en" ng-app="triviaApp">
3 <head>
4   <meta charset="utf-8" />
5   <meta http-equiv="x-ua-compatible" content="ie=edge">
6   <link rel="icon" href="https://www.triviaquestionsnow.com/favicon.ico" type="image/x-icon">
7   <meta name="viewport" content="width=device-width, initial-scale=1.0" />
8   <meta name="description" content="Want to put your sports knowledge to the test? triviaquestionsnow is the right place for you. Work your bra">
9   <meta name="keywords" content="Sports Trivia, Sports Trivia questions">
10
11  <link rel="apple-touch-icon" sizes="180x180" href="/apple-touch-icon.png">
12  <link rel="icon" type="image/png" sizes="32x32" href="/favicon-32x32.png">
13  <link rel="icon" type="image/png" sizes="16x16" href="/favicon-16x16.png">
14  <link rel="manifest" href="/site.webmanifest">
15  <link rel="mask-icon" href="/safari-pinned-tab.svg" color="#5bbad5">
16  <meta name="msapplication-TileColor" content="#da532c">
17  <meta name="theme-color" content="#ffffff">
18
19  <link rel="canonical" href="https://www.triviaquestionsnow.com/for/sports-trivia" />
20
21  <meta property="fb:app_id" content="329788460996850" />
22  <meta property="og:url" content="https://www.triviaquestionsnow.com/for/sports-trivia" />
23  <meta property="og:type" content="website" />
24    <meta property="og:title" content="Sports Trivia Questions and Answers" />
25    <meta property="og:description" content="Want to put your sports knowledge to the test? triviaquestionsnow is the right place for you." />
26    <meta property="og:image" content="https://www.triviaquestionsnow.com/trivia-questions-and-answers.jpg" />
27  <title>Sports Trivia Questions and Answers | TQN</title>
28  <!-- Fonts -->
29  <link rel="stylesheet" href="https://cdnjs.cloudflare.com/ajax/libs/font-awesome/4.5.0/css/font-awesome.min.css" integrity="sha384-XdYbMnZ/Q">
30  <link href="https://fonts.googleapis.com/css?family=Source+Sans+Pro:400,700,300" rel='stylesheet' type='text/css'>
31  <link href="https://cdnjs.cloudflare.com/ajax/libs/foundicons/3.0.0/foundation-icons.css" rel="stylesheet">
32  <!-- Styles -->
33  <link href="https://www.triviaquestionsnow.com/css/all.css?v=1" rel="stylesheet">
```

## `html_children()` and `html_node()`

Our tree has two children: `head` and `body`

```
html_obj %>% html_children()
```

```
## {xml_nodeset (2)}
## [1] <head>\n<meta http-equiv="Content-Type" content="text/html; charset=1
## [2] <body>\n      <div class="title-bar" data-responsive-toggle="nav" data-
```

Again notice the object returned might not be familiar  
("xml\_nodeset")

And each of the children has children of its own:

```
html_obj %>% html_node("body") %>% html_children()
```

```
## {xml_nodeset (6)}
## [1] <div class="title-bar" data-responsive-toggle="nav" data-hide-for="m
## [2] <div id="nav" class="top-bar">\n      <div class="row">\n
## [3] <div class="wrap bg-grey-light t-pad-20 b-pad-20">\n      <div
## [4] <script src="https://www.triviaquestionsnow.com/js/all.js?v=1"></scr
```

# html\_nodes()

Usually we'd figure out a rule and want a list of all relevant nodes:

```
html_obj %>% html_nodes("img")
```

```
## {xml_nodeset (8)}
## [1] % html_nodes("a")
```

```
## {xml_nodeset (44)}
## [1] <a href="/">\n                                \n## [3] <a href="https://www.triviaquestionsnow.com/easy-trivia-questions">\n## [4] <a href="https://www.triviaquestionsnow.com/for/sports-trivia">Sports\n## [5] <a href="https://www.triviaquestionsnow.com/for/music-trivia">Music
```

# html\_attrs()

Getting a specific attribute from those nodes:

```
html_obj %>% html_nodes("img") %>% html_attr("src")
```

```
## [1] "https://www.triviaquestionsnow.com/img/trivia-questions.png"
## [2] "https://www.triviaquestionsnow.com/img/trivia-questions.png"
## [3] "https://www.triviaquestionsnow.com/img/category/360x130/-category-1"
## [4] "https://www.triviaquestionsnow.com/img/category/360x130/-category-1"
## [5] "https://www.triviaquestionsnow.com/img/category/360x130/-category-1"
## [6] "https://www.triviaquestionsnow.com/img/category/360x130/apologetics"
## [7] "https://www.triviaquestionsnow.com/img/category/360x130/-category-1"
## [8] "https://www.triviaquestionsnow.com/img/category/360x130/-category-2"
```

```
html_obj %>% html_nodes("a") %>% html_attr("href")
```

```
## [1] "/"
## [2] "https://www.triviaquestionsnow.com"
## [3] "https://www.triviaquestionsnow.com/easy-trivia-questions"
## [4] "https://www.triviaquestionsnow.com/for/sports-trivia"
## [5] "https://www.triviaquestionsnow.com/for/music-trivia"
## [6] "https://www.triviaquestionsnow.com/for/math-trivia"
## [7] "https://www.triviaquestionsnow.com/categories"
```

# html\_text()

Getting the text from whatever set of elements we have:

```
html_obj %>% html_nodes("a") %>% html_text()
```

```
## [1] "\n"
## [2] "\n"
## [3] "Easy Trivia"
## [4] "Sports Trivia"
## [5] "Music Trivia"
## [6] "Math Trivia"
## [7] "Categories"
## [8] "All Trivia"
## [9] "\n"
## [10] "Show Answer"
## [11] "\n"
## [12] "Show Answer"
## [13] "\n"
## [14] "Show Answer"
## [15] "\n"
## [16] "Show Answer"
## [17] "\n"
## [18] "Show Answer"
## [19] "\n"
## [20] "Show Answer"
```

Which Country won the 2015 Davis Cup?\n

Which former NBA player had a starring role in t

Which Super Bowl had the most points ever scored

Who knocked Roger Federer out of the quarterfinals?

Which country was the first to win the Fed Cup, ?

In the sport of tennis, the acronym ATP stands for

# How to get to those questions? Option 1

Look at the page source, get some identifier yourself (class, ID, link)

Who was the only person to have won a Super Bowl as a player, as an assistant coach and as a head coach?

VeryHard

Show Answer

```
<div class="question callout" ng-controller="QuestionController as question">
<input type="hidden" ng-model="question.data.id" ng-init="question.data.id='1443'">
<span class="float-right light-grey bold l-cush-10">VeryHard</span>
<h3 class="fs-1 bold">
  <a href="https://www.triviaquestionsnow.com/question/person-to-win-super-bowl-as-player-as-assistant-coach-and-a
  Who was the only person to have won a Super Bowl as a player, as an assistant coach and as a head coach?
  </a>
</h3>
<div class="t-pad-10">
  <a href="#" class="click-to-show bold" ng-click="question.clickShow($event)">Show Answer</a>
  <div class="clearfix"></div>
  <p class="answer hide l-cush-0 t-cush-10" ng-class="{ 'answer': true, 'hide': !(question.showAnswer) }">
    <input type="hidden" name="question_position_1443" value="1">
    <input type="hidden" ng-model="question.data.id" ng-init="question.data.id='1443'">
    <input type="hidden" ng-model="question.data.answer" ng-init="question.data.answer='mike-ditka'">
    <strong>Answer: </strong><em>Mike Ditka</em>
    <br />
    <span class="dark-grey"><!--<strong>More information: </strong>-->Ditka won as a player with the Cowboys
  </p>
</div>
</div>
```

```
html_obj %>% html_nodes(".question") %>% .[[1]]
```

```
## {html_node}
## <div class="question callout" ng-controller="QuestionController as quest
## [1] <input type="hidden" ng-model="question.data.id" ng-init="question.d
## [2] <span class="float-right light-grey bold l-cush-10">Easy</span>
## [3] <h3 class="fs-1 bold">\n
## [4] <div class="t-pad-10">\n
```

## After some trial and error...

```
html_obj %>%
  html_nodes(".question") %>%
  html_nodes(".fs-1") %>%
  html_text() %>%
  str_trim()
```

```
## [1] "Which Country won the 2015 Davis Cup?"
## [2] "Which former NBA player had a starring role in the 1980 comedy movie"
## [3] "Which Super Bowl had the most points ever scored?"
## [4] "Who knocked Roger Federer out of the quarterfinals of the men's singles tournament at the 2012 U.S. Open?"
## [5] "Which country was the first to win the Fed Cup, the Davis Cup, and the World Cup in the same year?"
## [6] "In the sport of tennis, the acronym ATP stands for what?"
## [7] "The final of the men's doubles of the 2012 U.S. Open was won by which team?"
## [8] "Which of the following is not a possible outcome in an attempt to break a record in tennis?"
## [9] "\"Reggie's Prayer\" is a film starring with former NFL player?"
## [10] "The ground on which the game of golf is played is known as?"
```

# How to get to those questions? Option 2

SelectorGadget!



# From here it's a function fest!

```
extract_questions_and_answers_from_page <- function(url) {  
  html_obj <- read_html(url)  
  levels <- html_obj %>% html_nodes(".question") %>%  
    html_nodes(".l-cush-10") %>% html_text()  
  questions <- html_obj %>% html_nodes(".question") %>%  
    html_nodes(".fs-1") %>% html_text() %>% str_trim()  
  answers <- html_obj %>% html_nodes(".question") %>%  
    html_nodes(".answer") %>% html_text() %>%  
    str_extract(., "Answer:.*") %>% str_replace("Answer: ", "")  
  tibble(level = levels, question = questions, answer = answers)  
}  
  
extract_questions_and_answers_from_page(url)
```

```
## # A tibble: 10 × 3  
##   level    question  
##   <chr>    <chr>  
## 1 Medium   Which tennis player won the 2016 Laurens World Sports Award?  
## 2 Easy     How many teams are there currently in the NFL?  
## 3 Hard     The Doak Walker Award is given annually to the top college p.  
## 4 VeryHard Who was the first foreign-born NBA player to be selected num.  
## 5 Medium   Which NBA player once costarred alongside Jean-Claude Van Da.  
## 6 Easy     Who was the winner of the men's singles of the 2006 Tennis Ma.  
## 7 VeryHard When was the first year the three-point shot was introduced.  
## 8 Medium   Which player was the MVP of the 2014 NBA Final?
```

# Pagination

```
create_page_url <- function(topic, page_num) {  
  str_c("https://www.triviaquestionsnow.com/for/", topic, "-trivia")  
}  
  
extract_multiple_pages_single_topic <- function(topic, n = 5) {  
  cat(topic, "\n")  
  res <- map_dfr(  
    1:n,  
    function(i) {  
      cat(" ", i)  
      url <- create_page_url(topic, i)  
      extract_questions_and_answers_from_page(url)  
    }  
  )  
  res$topic <- topic  
  cat("\n")  
  res  
}
```

```
extract_multiple_pages_single_topic("sports")
```

```
## # A tibble: 50 × 4
##   level    question
##   <chr>    <chr>
## 1 VeryHard "In 1948, which NBA basketball team did the Harlem Glo... Minne...
## 2 Medium   "The Jacksonville Jaguars and Carolina Panthers entere... 1995
## 3 Hard     "An automatic progression by a player to the next stag... Bye
## 4 VeryHard "In 2016, Giants' wide receiver Odell Beckham, Jr. app... Code
## 5 Hard     "Which NBA player broke the record for most points sco... Jerem...
## 6 Medium   "Before relocating to Foxborough, Massachusetts, what ... Boston
## 7 Easy     "What is the term for the historic jerseys today worn ... Throw...
## 8 Medium   "Who served as the starting center of the Golden State... Andre...
## 9 Easy     "In what year was the 4 minute mile achieved?"          1954
## 10 Hard    "Who was the first tennis player to complete a \"Grand... Don ...
## # ... with 40 more rows
```

# Magic!

```
topics <- c("sports", "kids", "science", "bible", "food-drink", "geography", "history", "video-games")  
  
df_all <- map_dfr(  
  topics,  
  extract_multiple_pages_single_topic  
)  
  
df_all %>% count(topic)  
  
## # A tibble: 8 × 2  
##   topic             n  
##   <chr>            <int>  
## 1 bible              50  
## 2 food-drink         50  
## 3 geography          50  
## 4 history             50  
## 5 kids                50  
## 6 science              50  
## 7 sports               50  
## 8 video-games         50
```

# BeautifulSoup

APPLICATIONS



OF DATA SCIENCE

# Almost always start with

```
import requests
from bs4 import BeautifulSoup

html_obj = requests.get('https://en.wikipedia.org/wiki/List_of_The
soup = BeautifulSoup(html_obj.content, 'html.parser')
type(soup)

## <class 'bs4.BeautifulSoup'>
```

This object has all sorts of attributes and methods:

```
soup.get_text()
soup.prettify()
soup.attrs
soup.children
soup.title
```

## **find() a tag, find\_all()**

```
link_objs = soup.find_all('a', href=True)
type(link_objs)

## <class 'bs4.element.ResultSet'>

type(link_objs[3])

## <class 'bs4.element.Tag'>

link_objs[3].text

## 'Current events'

link_objs[3].attrs

## {'href': '/wiki/Portal:Current_events', 'title': 'Articles related to cu
```

See the actual link in the [page](#).

# Getting that table

```
table = soup.find('table', attrs={'class':'wikitable'})  
table_body = table.find('tbody')  
  
rows = table_body.find_all('tr')  
print(len(rows))
```

## 155

```
print(rows[0])
```

```
## <tr>  
## <th rowspan="2">Installment  
## </th>  
## <th rowspan="2">Housewives  
## </th>  
## <th rowspan="2">First season<br/>starred  
## </th>  
## <th rowspan="2">Last season<br/>starred  
## </th>  
## <th colspan="4">Number of seasons  
## </th></tr>
```

# Getting a Housewife name

```
<tr>
<td><span data-sort-value="De La Rosa, Jo&#160;!">Jo De La Rosa</span>
</td>
<td>1
</td>
<td>2
</td>
<td>2
</td>
<td>0
</td>
<td>2
</td></tr>
<tr>
<td><span data-sort-value="Gunvalson, Vicki&#160;!"><a href="/wiki/Vicki_Gunvalson" title="Vicki Gunvalson">Vicki Gunvalson</a></span>
</td>
<td>1
</td>
<td>13
```

```
import re

print(rows[3].find('span', attrs = { 'data-sort-value': re.compile

## <span data-sort-value="De La Rosa, Jo !">Jo De La Rosa</span>
```

# Getting only HWives with Wiki pages

```
housewives_with_links = []
for row in rows:
    housewife = row.find('span',
        attrs = {'data-sort-value': re.compile(r'.*')})
    if housewife is not None:
        link = housewife.find('a')
        if link is not None:
            housewives_with_links.append((housewife.text, link['href']))

import pandas as pd

h_df = pd.DataFrame(housewives_with_links, columns=['name', 'link'])

h_df.head()
```

```
##                                     name                               link
## 0          Vicki Gunvalson /wiki/Vicki_Gunvalson
## 1          Jeana Keough  /wiki/Jeana_Keough
## 2          Tamra Judge  /wiki/Tamra_Judge
## 3          Heather Dubrow /wiki/Heather_Dubrow
## 4 Shannon Storms Beador /wiki/Shannon_Storms_Beador
```

# (Though if your table is simple, try:)

```
l = pd.read_html(html_obj.text)  
l[0].head()
```

```
##      Installment      Housewives ... Number of seasons  
##      Installment      Housewives ...           Guest Ultimate Girls  
## 0  Orange County  Kimberly Bryant ...             3  
## 1  Orange County   Jo De La Rosa ...            2  
## 2  Orange County  Vicki Gunvalson ...           0  
## 3  Orange County    Jeana Keough ...            5  
## 4  Orange County    Lauri Peterson ...          1  
##  
## [5 rows x 8 columns]
```

# Following HWives Links

```
def get_housewife_img_ref(housewife_link):
    html_obj = requests.get('https://en.wikipedia.org' + housewife_link)
    soup = BeautifulSoup(html_obj.content, 'html.parser')
    infobox = soup.find('table', attrs = {'class': 'vcard'})
    if infobox is not None:
        img_obj = infobox.find('img', src=True)
        if img_obj is not None:
            return img_obj['src']
    return None

h_df['img_ref'] = h_df['link'].apply(get_housewife_img_ref)
h_df.dropna(inplace=True)

h_df.head()
```

```
##                                     name ...
## 0          Vicki Gunvalson ...
## 1          Luann de Lesseps ...
## 2          Bethenny Frankel ...
## 3  Kelly Killoren Bensimon ...
## 4          Carole Radziwill ...
##
## [5 rows x 3 columns]
```

# Downloading HWives Images

```
def make_img_filename(hf_name):
    return 'data/housewives/' + hf_name.lower().strip(',.-').replace(' ', '_')

def download_hw_img(hf_name, hf_img_ref):
    img_file = make_img_filename(hf_name)
    img_data = requests.get('http:' + hf_img_ref).content
    with open(img_file, 'wb') as handler:
        handler.write(img_data)

h_df.apply(lambda row: download_hw_img(row['name'], row['img_ref'],
                                         axis=1))
```