

APPLICATIONS



OF DATA SCIENCE

Visualizing in the Tidyverse

Applications of Data Science - Class 4

Giora Simchoni

gsimchoni@gmail.com and add #dsapps in subject

Stat. and OR Department, TAU

2021-02-09

APPLICATIONS



OF DATA SCIENCE

Welcome to ggplot2

Heavily inspired by datascienceinabox.org / @minebocek

APPLICATIONS



OF DATA SCIENCE

ggplot()

- **ggplot2** is tidyverse's data visualization package
- The **gg** in "ggplot2" stands for Grammar of Graphics
- It is inspired by the book [Grammar of Graphics](#) by Leland Wilkinson
- A grammar of graphics is a tool that enables us to concisely describe the components of a graphic



Remember the Force, Luke!

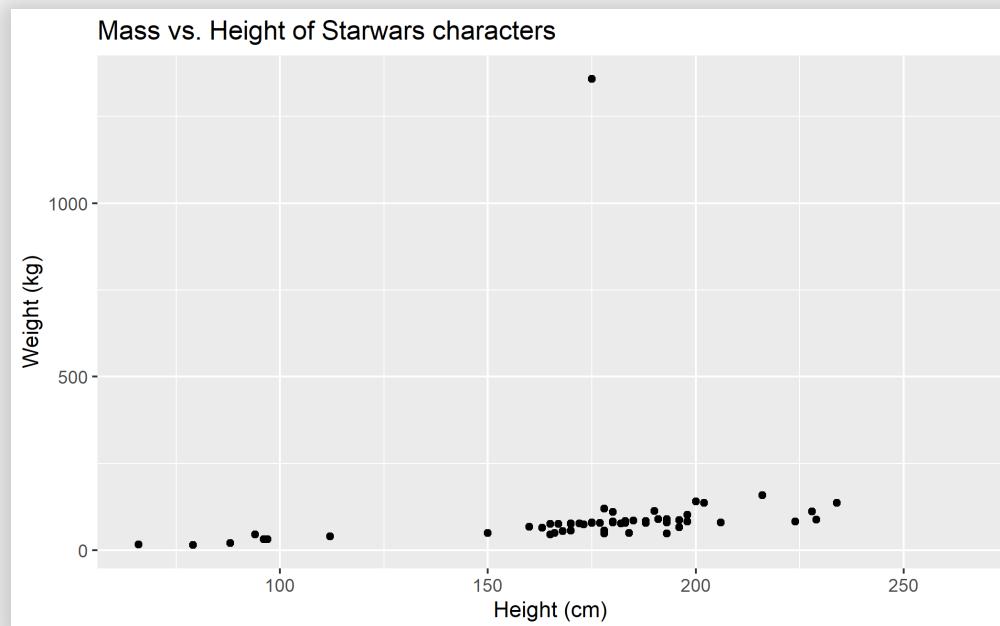
```
sw_tables <- read_rds("../data/sw_tables.rds")
characters <- sw_tables$characters
planets <- sw_tables$planets
films <- sw_tables$films

qglimpse(characters)
```

What are the functions doing the plotting? What is the dataset being plotted? Which variable is on the x-axis and which variable is on the y-axis? What does the warning mean?

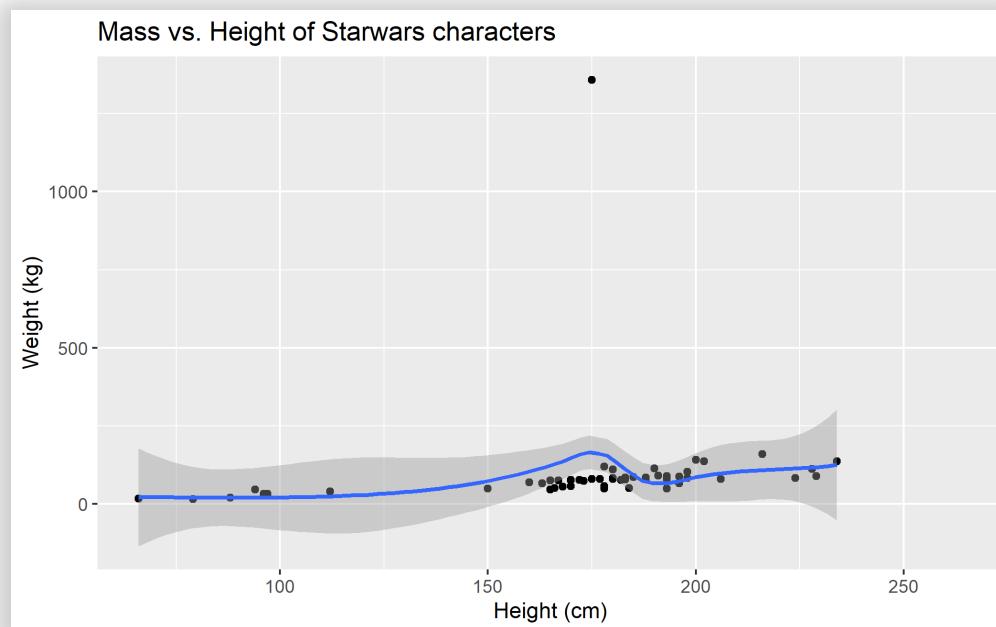
```
ggplot(data = characters, mapping = aes(x = height, y = mass)) +  
  geom_point() +  
  labs(title = "Mass vs. Height of Starwars characters",  
       x = "Height (cm)", y = "Weight (kg)")
```

`## Warning: Removed 36 rows containing missing values (geom_point).`



What does `geom_smooth()` do? What else changed between the previous plot and this one?

```
ggplot(data = characters, mapping = aes(x = height, y = mass)) +  
  geom_point() +  
  geom_smooth() +  
  labs(title = "Mass vs. Height of Starwars characters",  
       x = "Height (cm)", y = "Weight (kg)")
```



Who dis!

- `ggplot()` is the main function in `ggplot2` and plots are constructed in layers
- The structure of the code for plots can often be summarized as

```
ggplot +  
  geom_xxx
```

or, more precisely

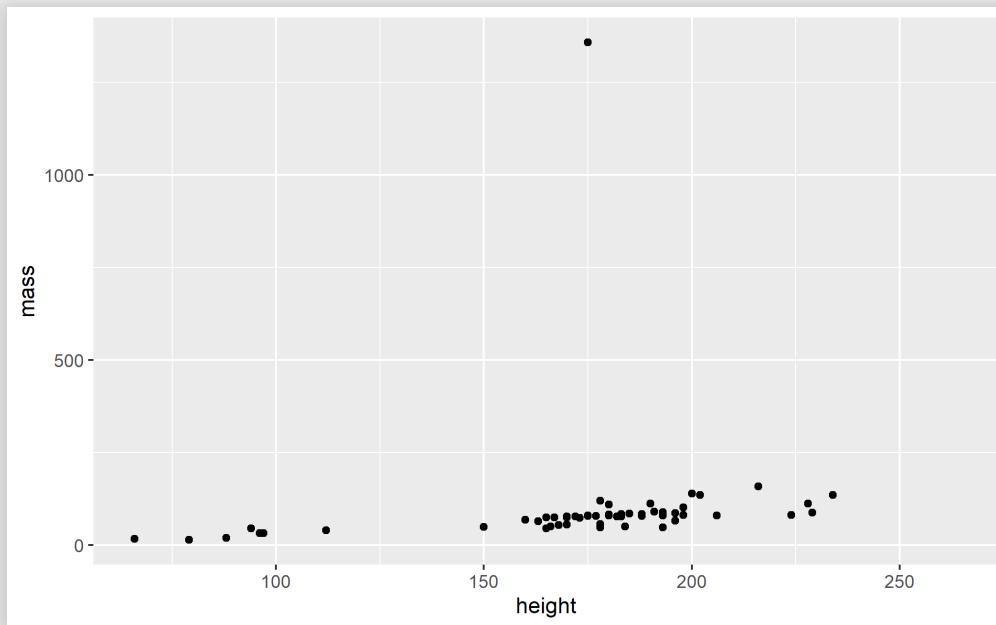
```
ggplot(data = [dataset], mapping = aes(x = [x-variable], y = [y-varaible]),  
       geom_xxx() +  
       other options)
```

- To use `ggplot2` functions, first load `tidyverse`

```
library(tidyverse)
```

Mass vs. Height

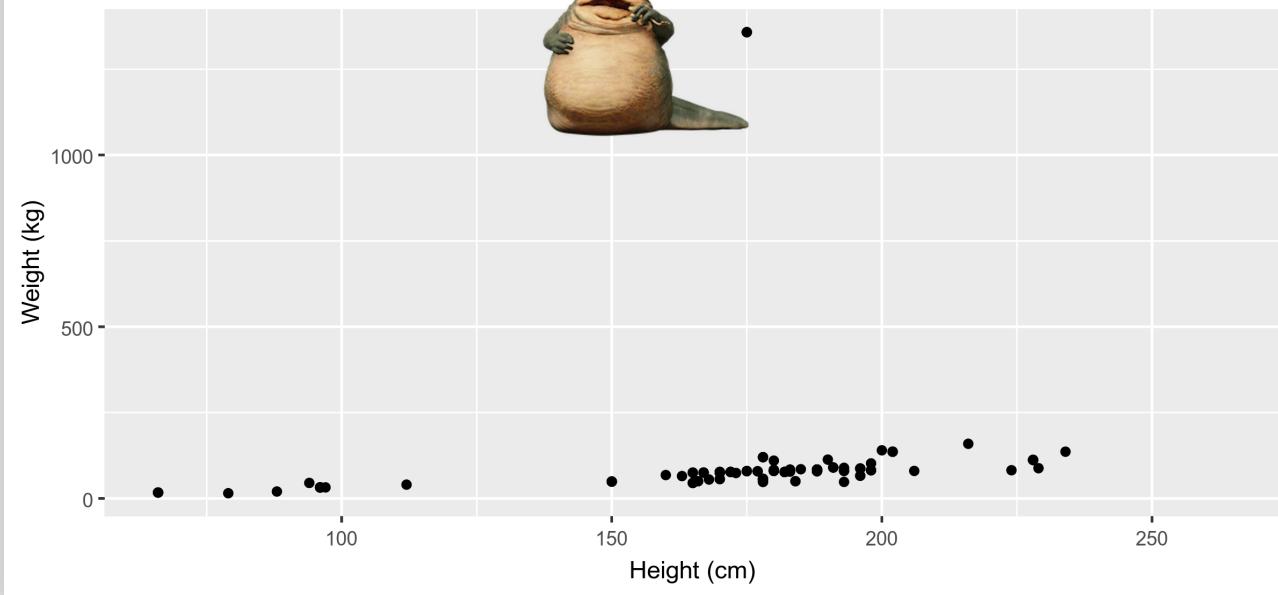
```
ggplot(characters, aes(height, mass)) + geom_point()
```



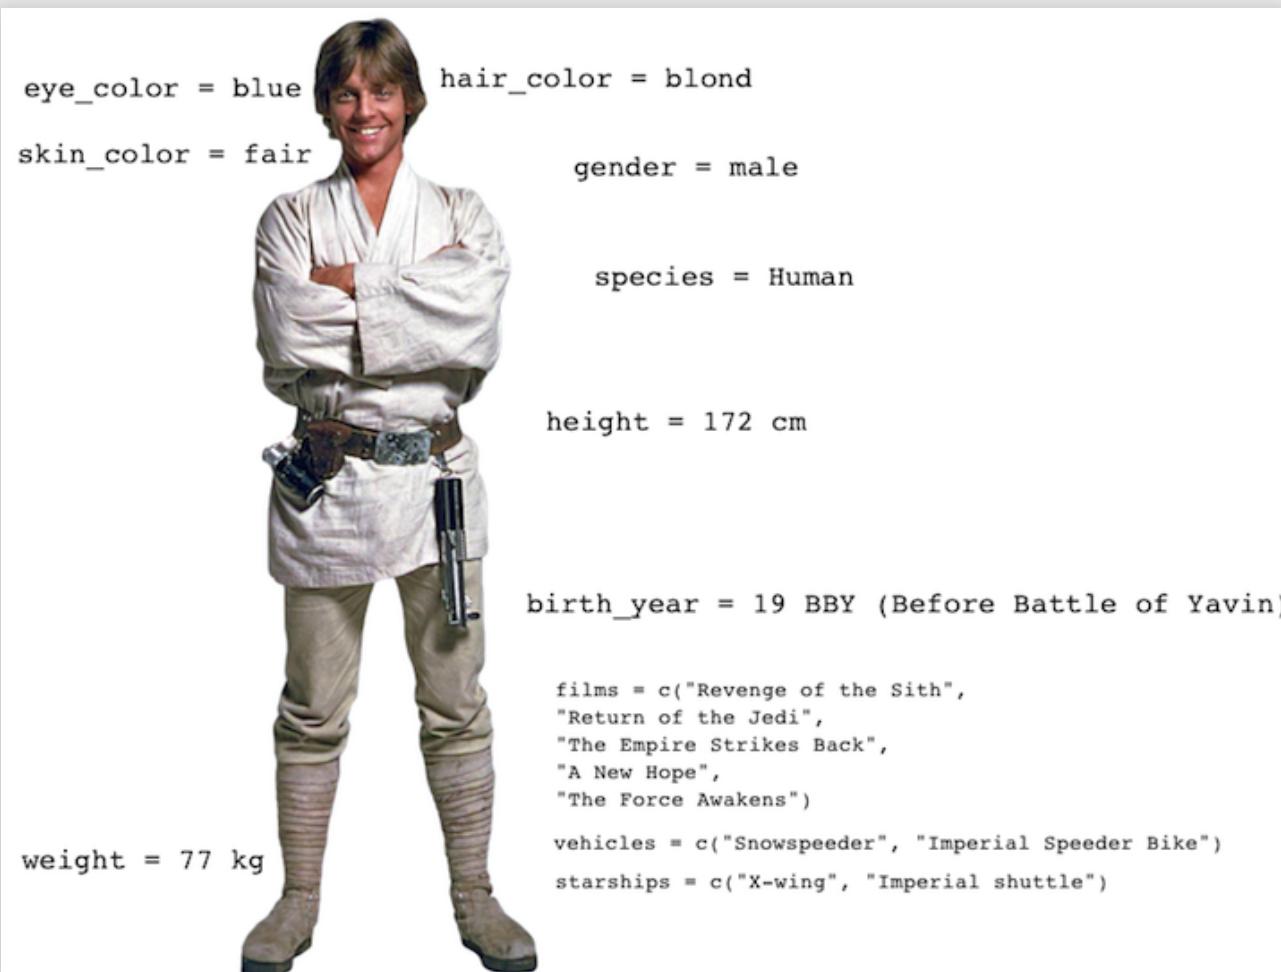
How would you describe this relationship? What other variables would help us understand data points that don't follow the overall trend? Who is the not so tall but really chubby character?

Jabba!

Mass vs. Height of Starwars characters



But seriously, Luke



Additional variables

We can map additional variables to various features of the plot:

- aesthetics
 - shape
 - colour
 - size
 - alpha (transparency)
- faceting: small multiples displaying different subsets

Aesthetics

APPLICATIONS



OF DATA SCIENCE

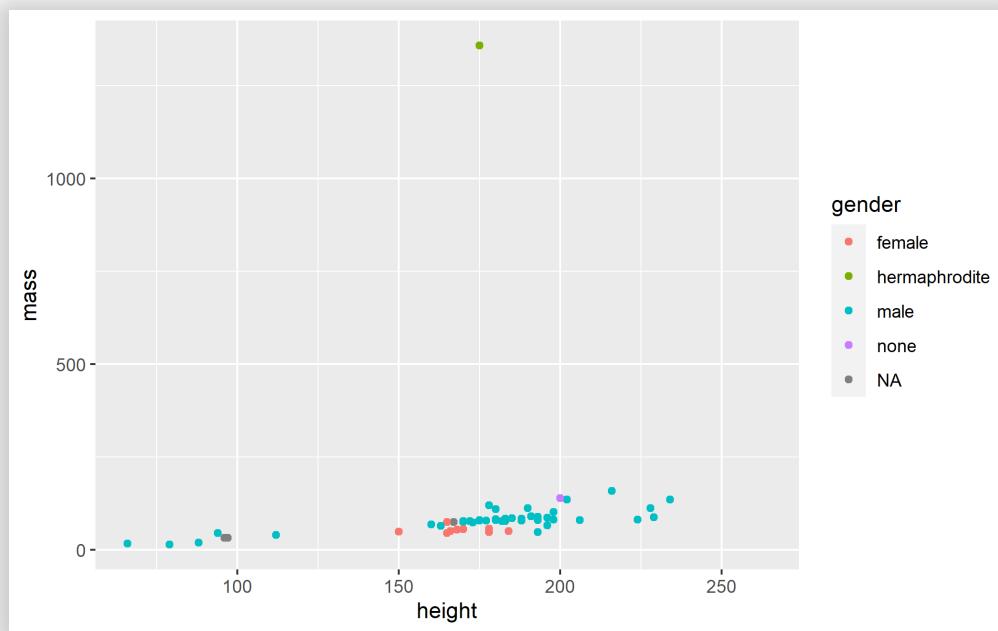
Aesthetics options

Visual characteristics of plotting characters that can be **mapped to a specific variable** in the data are

- color
- fill
- size
- shape
- alpha (transparency)

Mass vs. Height + Gender

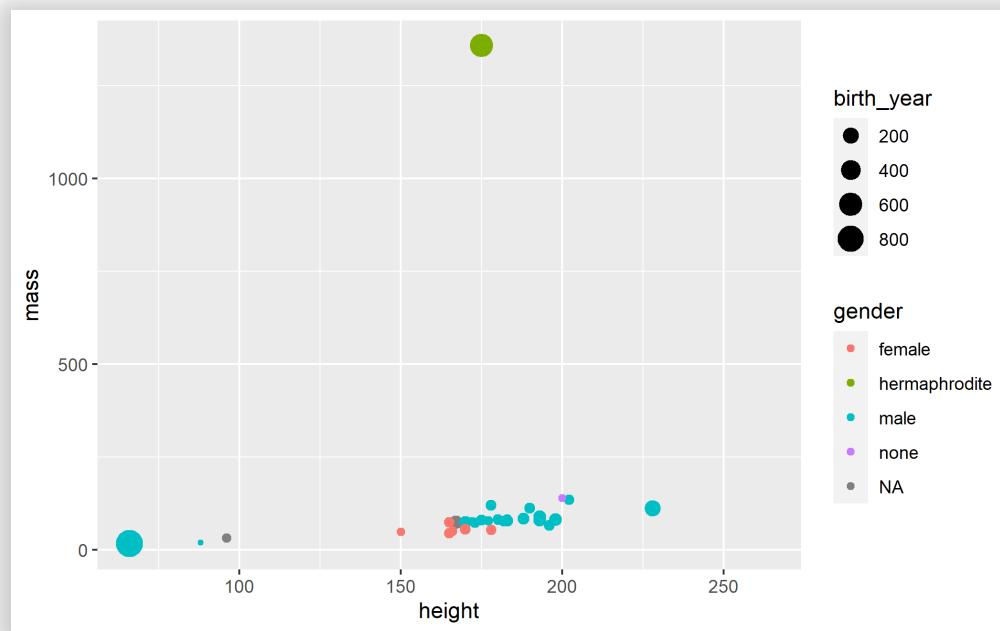
```
ggplot(data = characters,  
       mapping = aes(x = height, y = mass, color = gender)) +  
  geom_point()
```



Mass vs. Height + Gender + Birth Year

Let's map the size to birth_year:

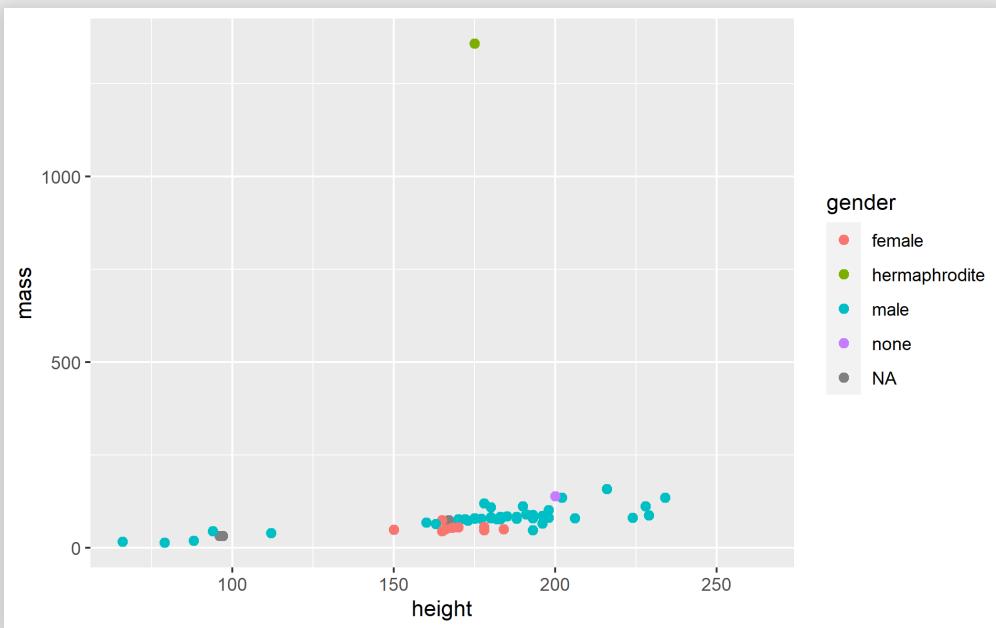
```
ggplot(data = characters,  
       mapping = aes(x = height, y = mass, color = gender,  
                      size = birth_year)) +  
  geom_point()
```



Mass vs. Height + Gender

Let's now increase the size of all points **not** based on the values of a variable in the data:

```
ggplot(data = characters,  
       mapping = aes(x = height, y = mass, color = gender)) +  
  geom_point(size = 2)
```



Aesthetics Summary

- Continuous variables are measured on a continuous scale
- Discrete variables are measured (or often counted) on a discrete scale

aesthetics	discrete	continuous
color	rainbow of colors	gradient
size	discrete steps	linear mapping between radius and value
shape	different shape for each	shouldn't (and doesn't) work

- Use aesthetics for mapping features of a plot to a variable, define the features in the geom for customization **not** mapped to a variable

Faceting

APPLICATIONS



OF DATA SCIENCE

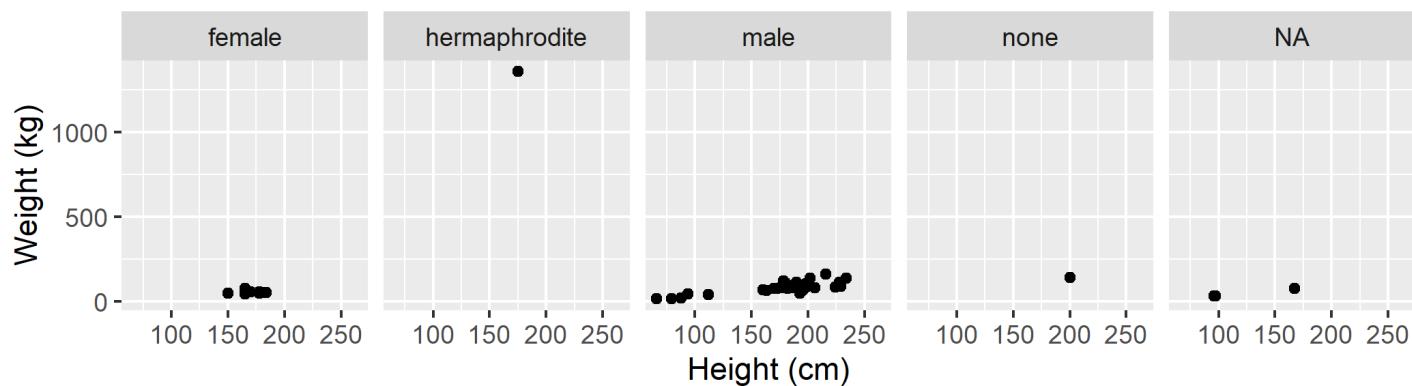
Faceting options

- Smaller plots that display different subsets of the data
- Useful for exploring conditional relationships and large data

```
ggplot(data = characters, mapping = aes(x = height, y = mass)) +  
  facet_grid(. ~ gender) +  
  geom_point() +  
  labs(title = "Mass vs. Height of Starwars characters",  
       subtitle = "Faceted by Gender",  
       x = "Height (cm)", y = "Weight (kg)")
```

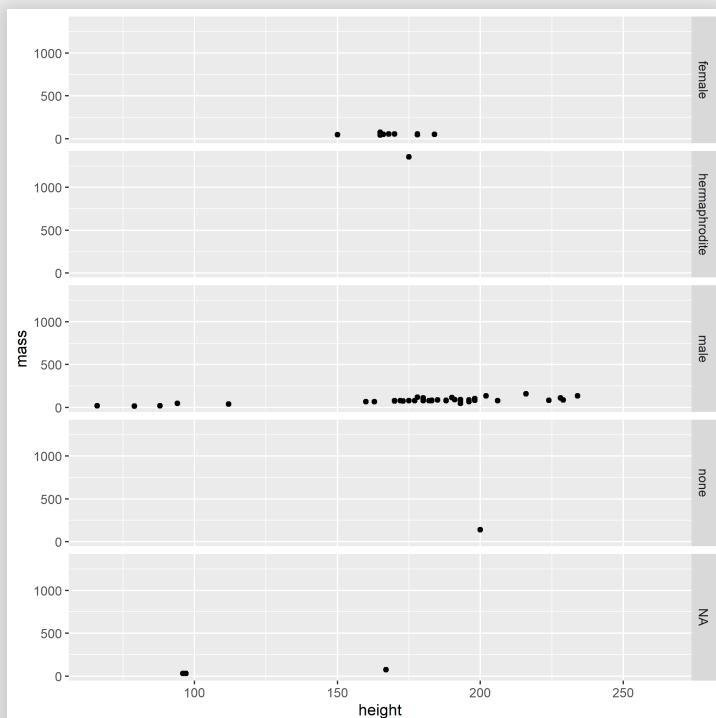
Mass vs. Height of Starwars characters

Faceted by Gender

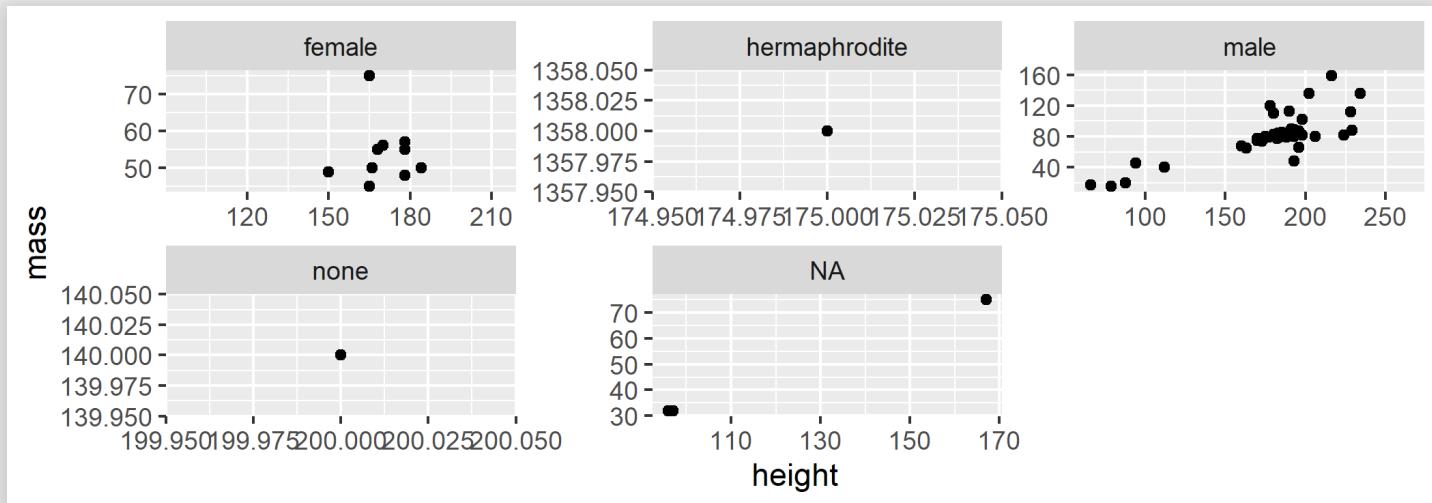


Many ways to facet

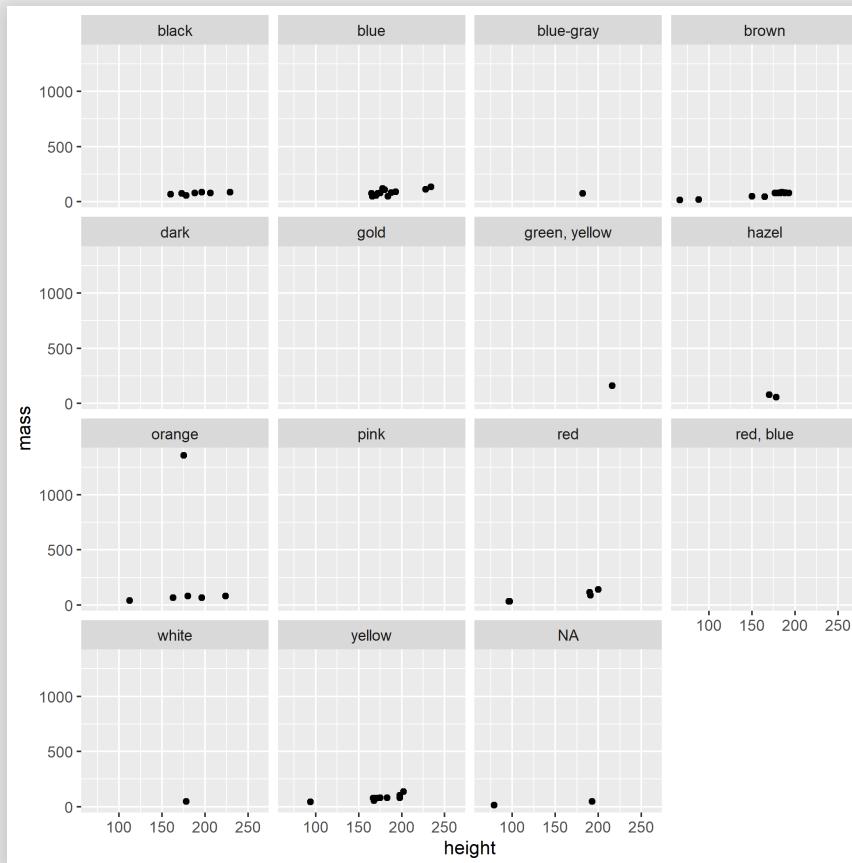
```
ggplot(data = characters, mapping = aes(x = height, y = mass)) +  
  geom_point() +  
  facet_grid(gender ~ .)
```



```
ggplot(data = characters, mapping = aes(x = height, y = mass)) +  
  geom_point() +  
  facet_wrap(. ~ gender, scales = "free")
```



```
ggplot(data = characters, mapping = aes(x = height, y = mass)) +  
  geom_point() +  
  facet_wrap(~ eye_color)
```



Facet summary

- `facet_grid()`:
 - 2d grid
 - `rows ~ cols`
 - use `.` for no split
- `facet_wrap()`: 1d ribbon wrapped into 2d

Take Control

APPLICATIONS



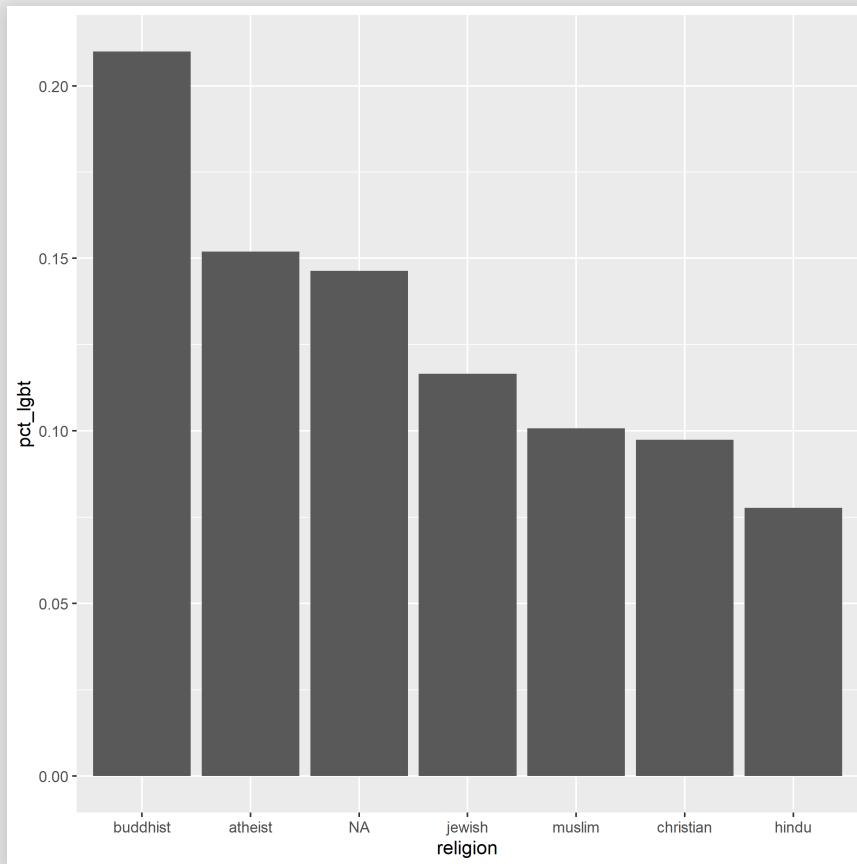
OF DATA SCIENCE

```
okcupid_lgbt_relig <- okcupid %>%
  group_by(religion2) %>%
  summarise(pct_lgbt = mean(orientation %in% c("gay", "bisexual")))
  mutate(religion = fct_reorder(religion2, desc(pct_lgbt)))

okcupid_lgbt_relig
```

```
## # A tibble: 7 x 3
##   religion2 pct_lgbt religion
##   <chr>        <dbl> <fct>
## 1 atheist      0.152 atheist
## 2 buddhist     0.210 buddhist
## 3 christian    0.0974 christian
## 4 hindu        0.0778 hindu
## 5 jewish       0.117 jewish
## 6 muslim        0.101 muslim
## 7 NA            0.146 NA
```

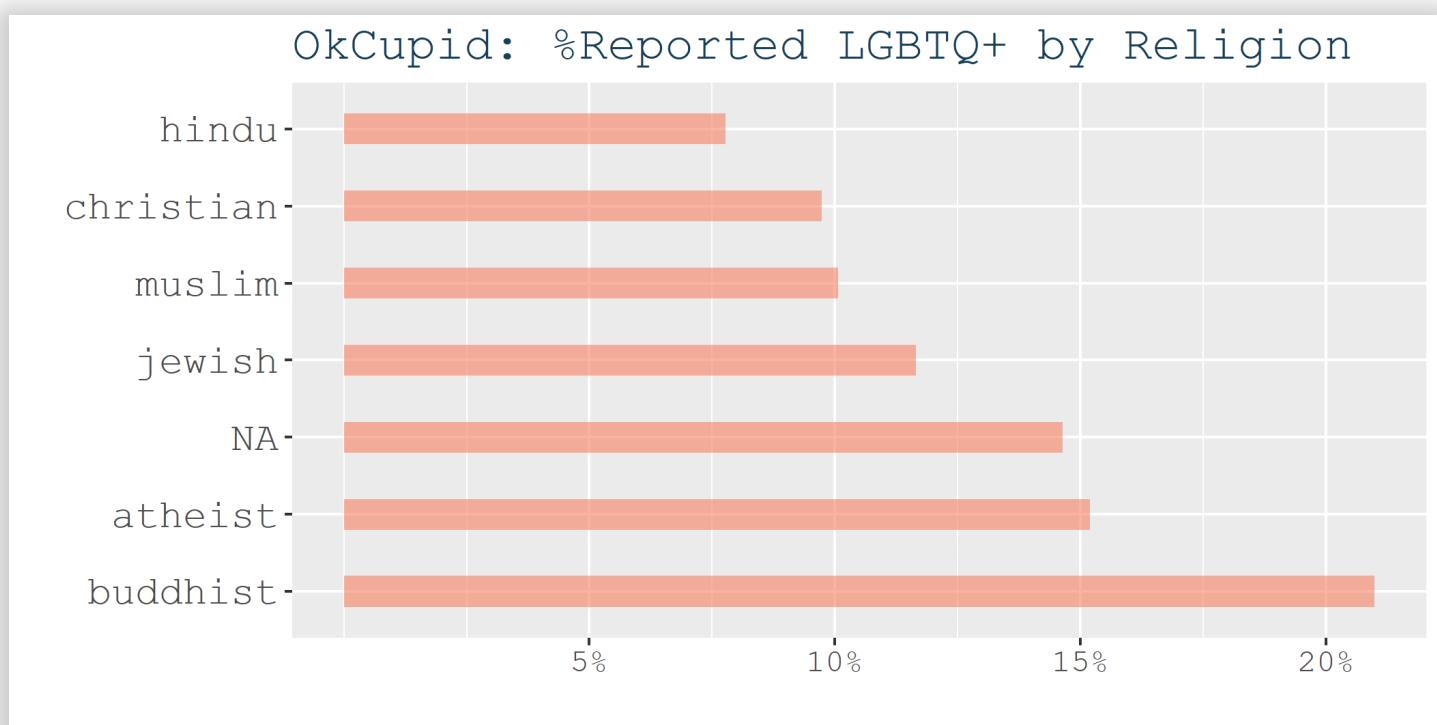
```
okcupid_lgbt_relig %>%  
  ggplot(aes( religion, pct_lgbt)) +  
  geom_bar(stat="identity")
```



Pimp my plot

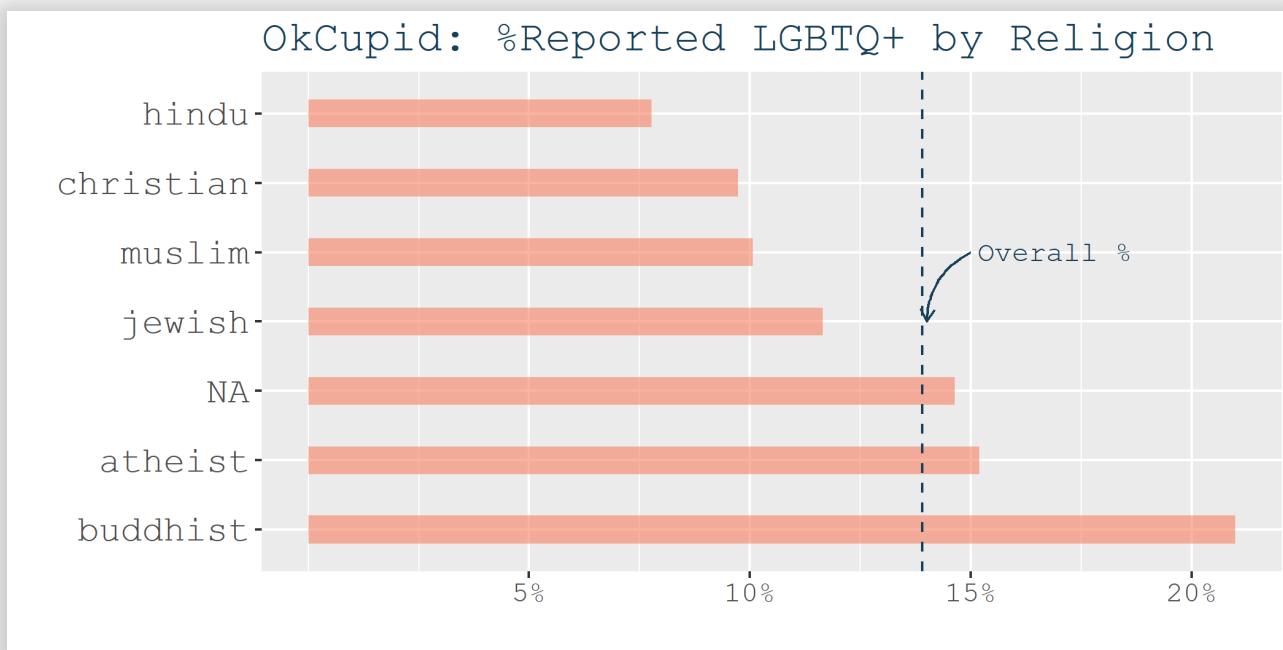
```
p <- ggplot(okcupid_lgbt_relig, aes(religion, pct_lgbt)) +
  geom_bar(stat="identity", fill="#f68060", alpha=.6, width=.4) +
  labs(x = "", y = "",
       title = "OkCupid: %Reported LGBTQ+ by Religion") +
  theme(text = element_text(size = 14, color = "#163f59", family =
    axis.text.y = element_text(size = 14)) +
    scale_y_continuous(breaks = seq(0.05, 0.2, 0.05),
                       labels=scales::percent_format(accuracy = 1)))
  coord_flip()
p
```

Pimp my plot



Annotate, Add lines

```
p + geom_hline(yintercept = 0.139, lty = 2, color = "#163f59") +  
  annotate("text", x = 5, y = 0.139 + 0.03, label = "Overall %",  
          color = "#163f59", family = "mono") +  
  annotate(geom = "curve", x = 5, y = 0.15, xend = 4, yend = 0.139,  
          curvature = .3, arrow = arrow(length = unit(2, "mm"))), color =
```



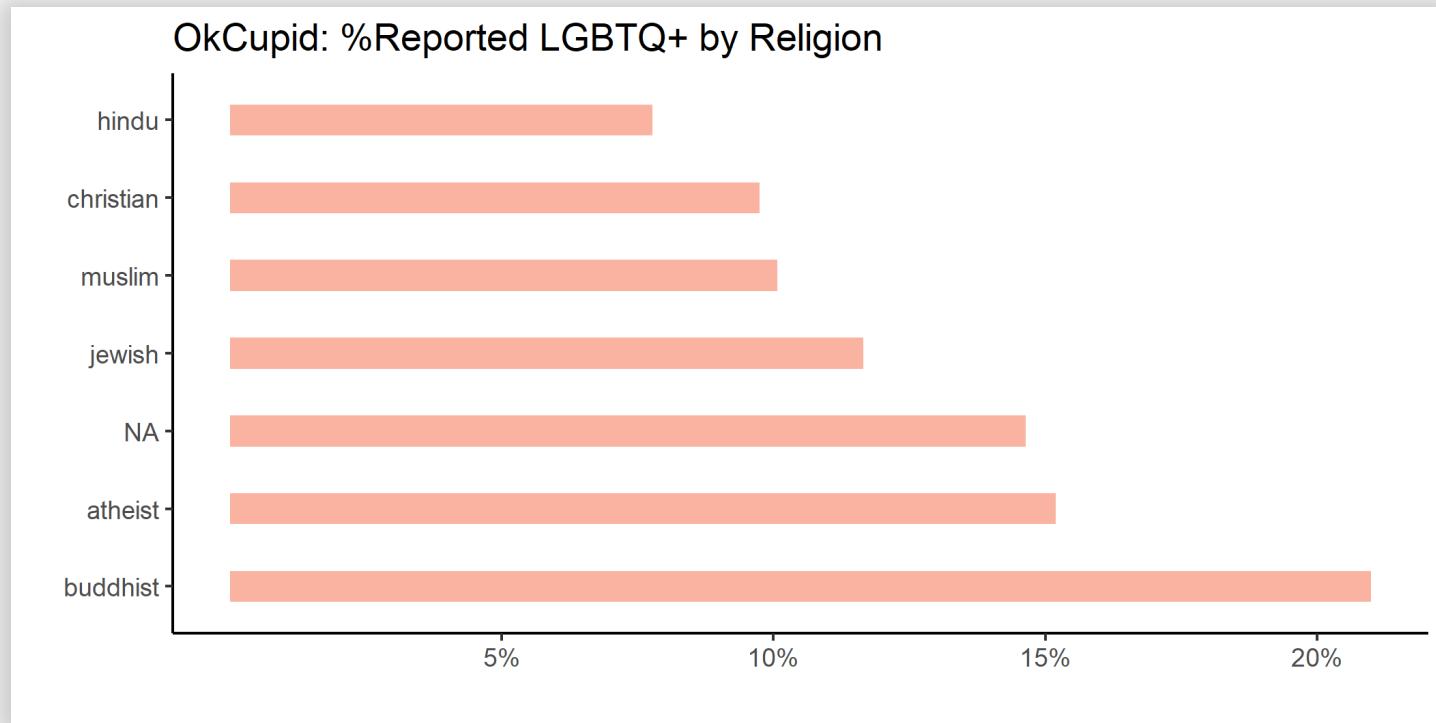
Themes

APPLICATIONS

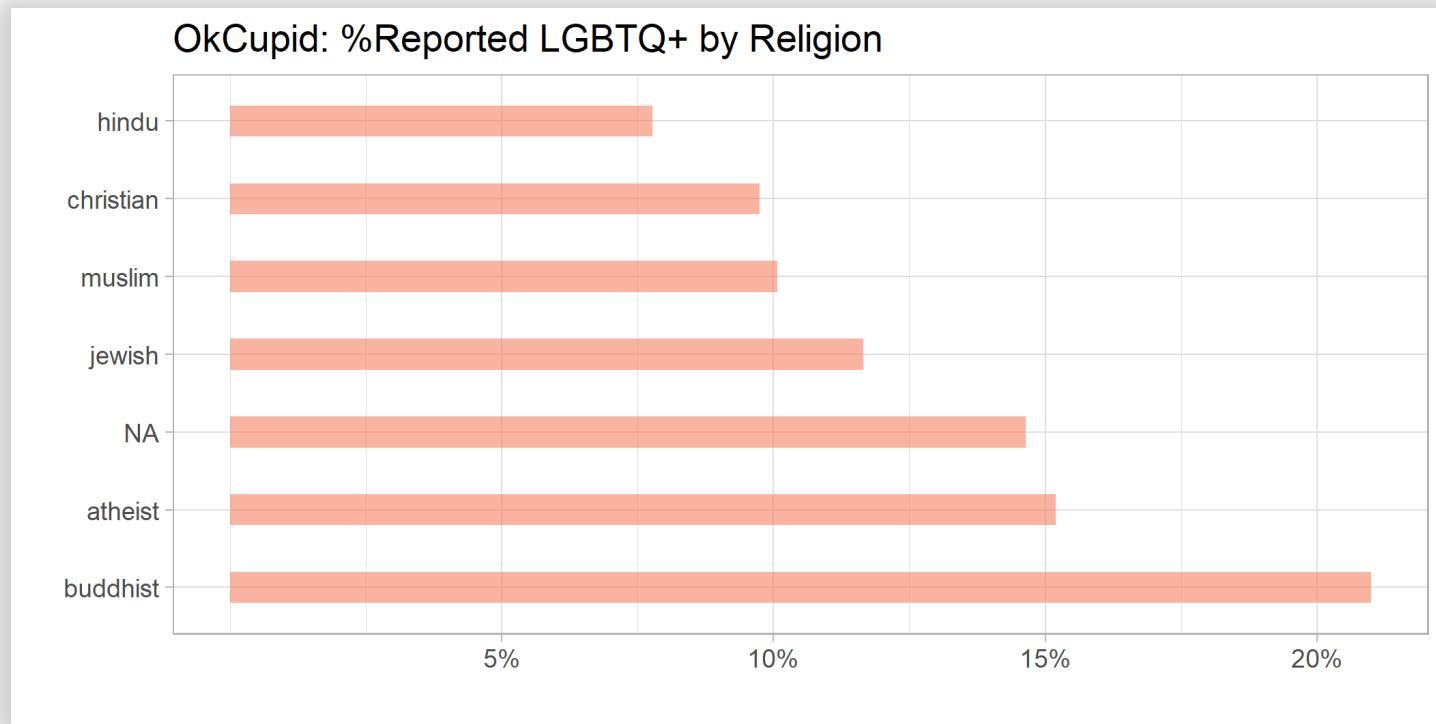


OF DATA SCIENCE

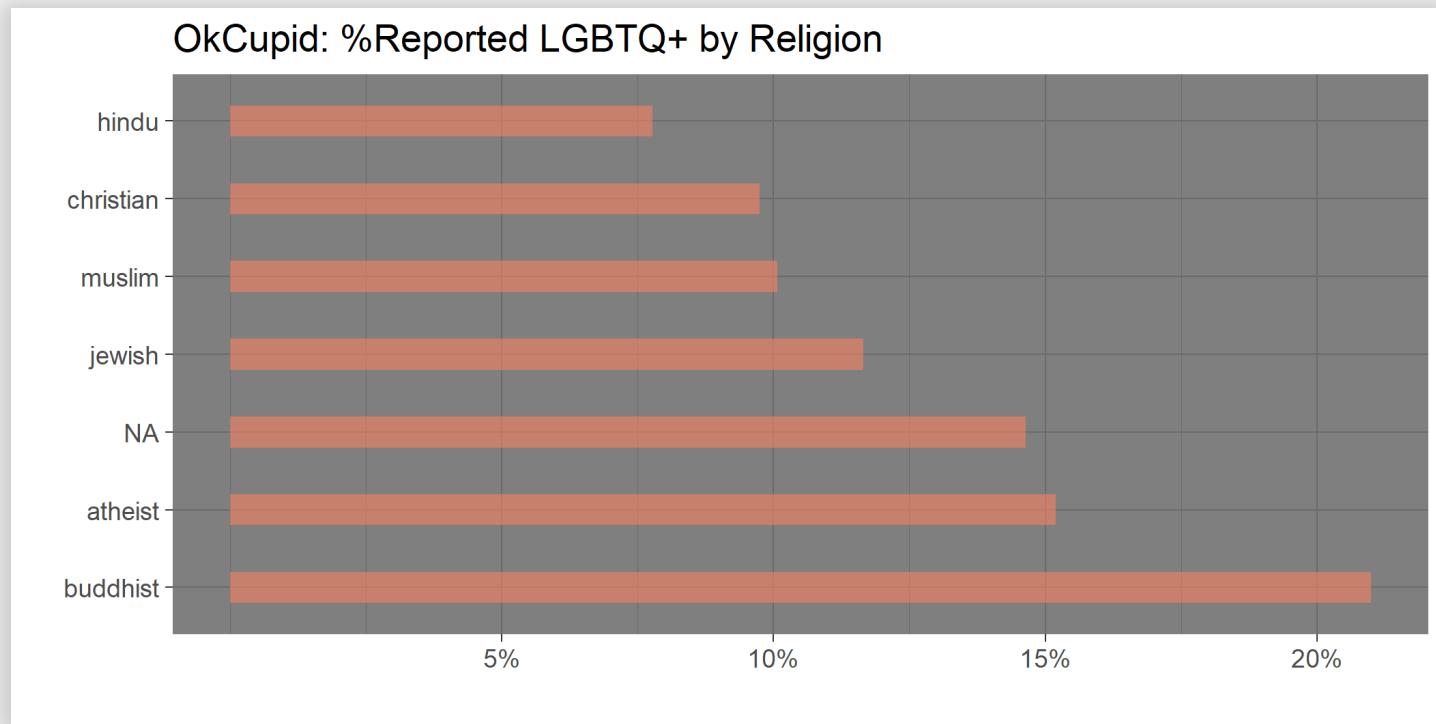
```
p + theme_classic()
```



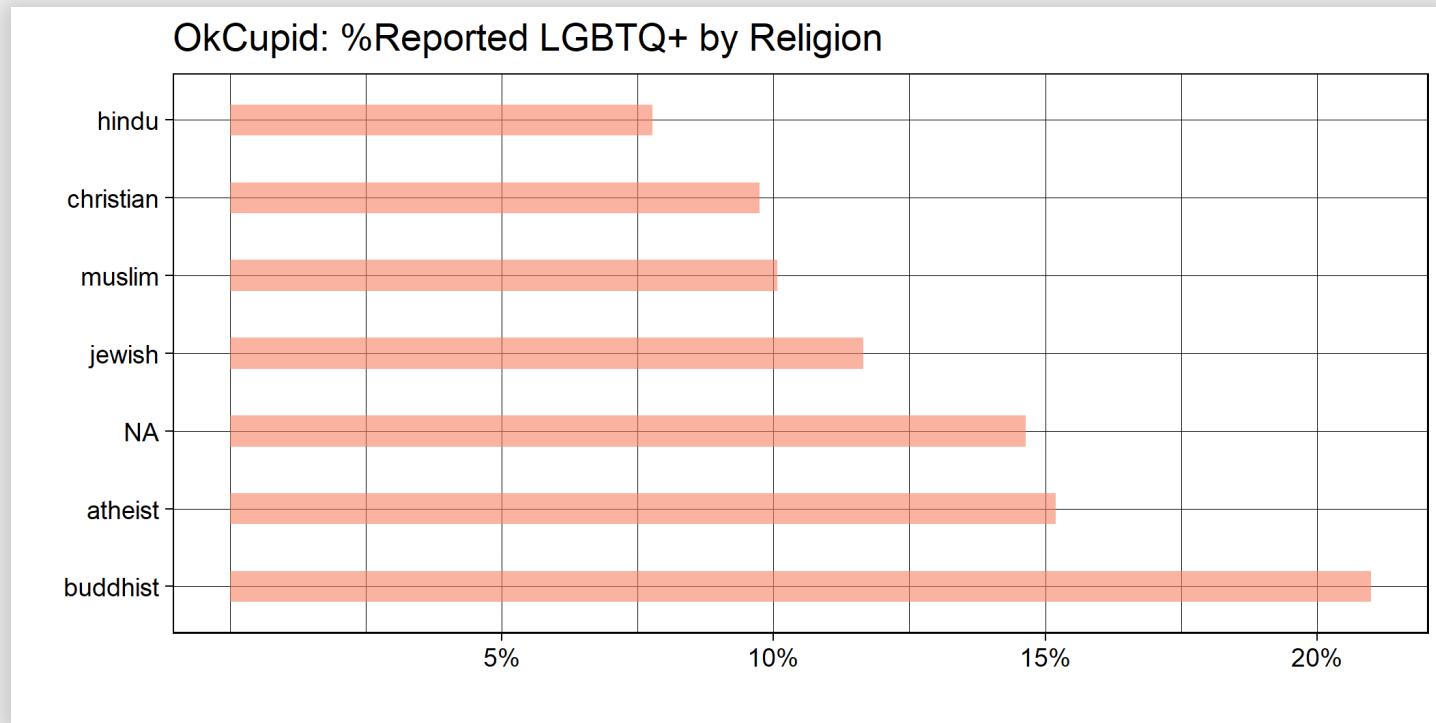
```
p + theme_light()
```



```
p + theme_dark()
```

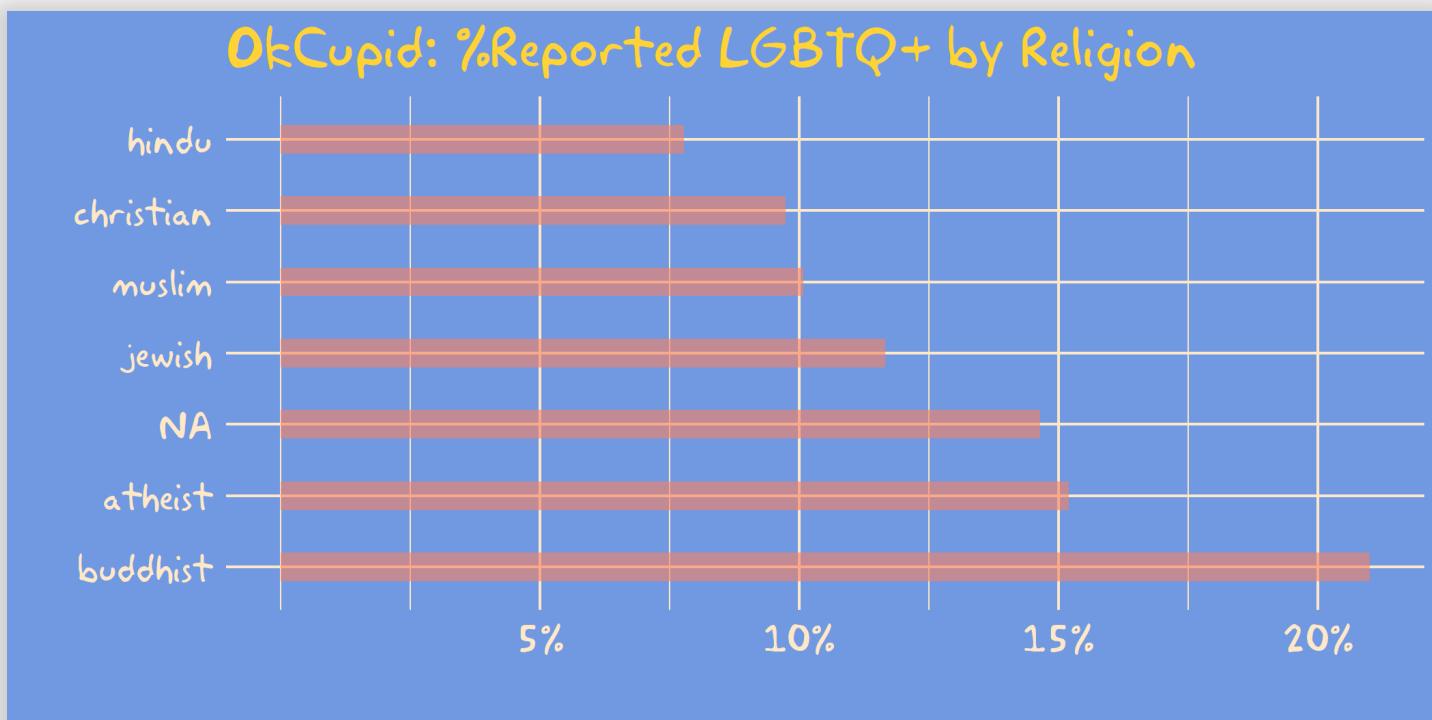


```
p + theme_linedraw()
```



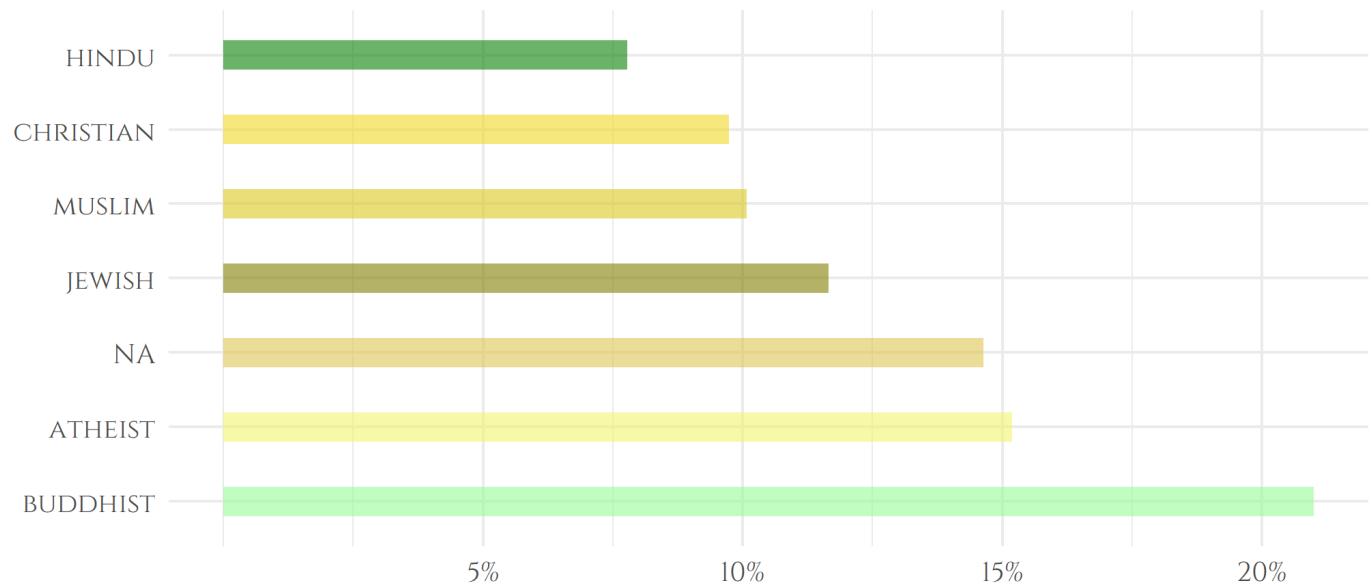
```
library(tvthemes)
# import_simpsons()
extrafont::loadfonts(device="win")

p + theme_simpsons(title.font = "Akbar",
                     text.font = "Akbar",
                     axis.text.size = 12)
```



```
ggplot(okcupid_lgbt_relig, aes(religion, pct_lgbt)) +  
  geom_bar(aes(fill = pct_lgbt), stat="identity", alpha=.6, width=1)  
  labs(x = "", y = "", title = "This was supposed to be GoT...") +  
  scale_y_continuous(breaks = seq(0.05, 0.2, 0.05), labels=scales::percent) +  
  coord_flip() +  
  scale_fill_westeros(palette = "Tyrell", type = "continuous") +  
  theme(text = element_text(family = "Cinzel", size = 12),  
        title = element_text(family = "Cinzel", size = 14)) +  
  guides(fill=guide_colorbar("Tyrell"))
```

THIS WAS SUPPOSED TO BE GOT...



Common Plots à la GG

APPLICATIONS



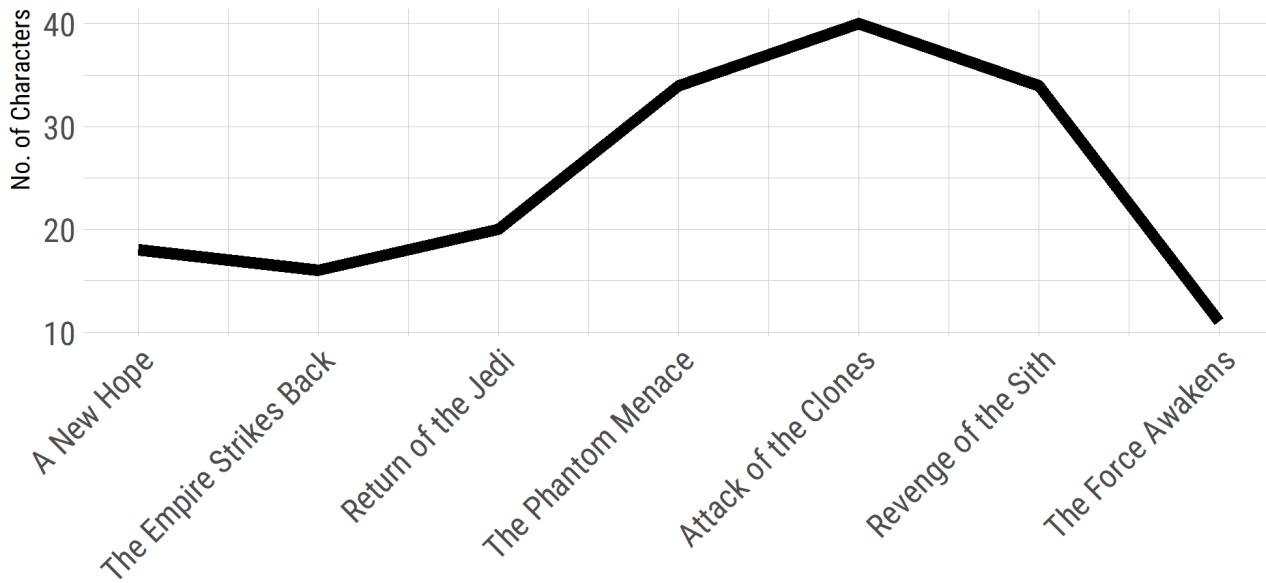
OF DATA SCIENCE

Line Plot

```
library(hrbrthemes)

characters %>%
  group_by(film_id) %>%
  summarise(n_char = n_distinct(character_id)) %>%
  ggplot(aes(film_id, n_char)) +
  geom_line(lwd = 2) +
  labs(x = "", y = "No. of Characters",
       title = "Star Wars Films No. of Unique Characters") +
  scale_x_continuous(breaks = 1:7, labels = films$title) +
  theme_ipsum(base_family = "Roboto Condensed") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

Star Wars Films No. of Unique Characters



Histogram

```
library(scales)

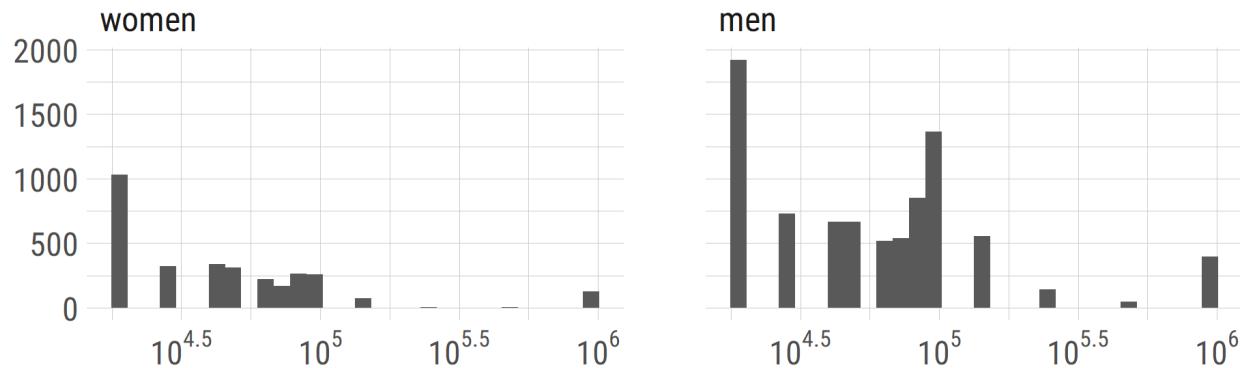
sex_names <- list(
  "m" = "men",
  "f" = "women"
)

sex_labeller <- function(variable, value){
  return(sex_names[value])
}

p2 <- okcupid %>%
  ggplot(aes(income)) +
  facet_wrap(. ~ sex, labeller = sex_labeller) +
  scale_x_log10(breaks = trans_breaks("log10", function(x) 10^x),
                 labels = trans_format("log10", math_format(10^.x)))
  labs(x = "", y = "", title = "OkCupid: Annual Income Dist. by Ge
  caption = "Income is presented as log10(Income)") +
  theme_ipsum(base_family = "Roboto Condensed")

p2 + geom_histogram()
```

OkCupid: Annual Income Dist. by Gender

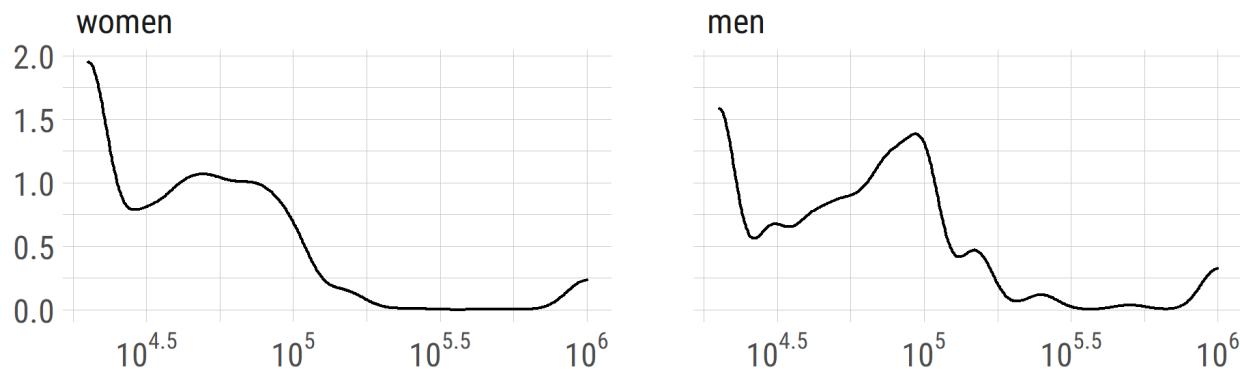


Income is presented as $\log_{10}(\text{Income})$

Density Plot

```
p2 + geom_density()
```

OkCupid: Annual Income Dist. by Gender



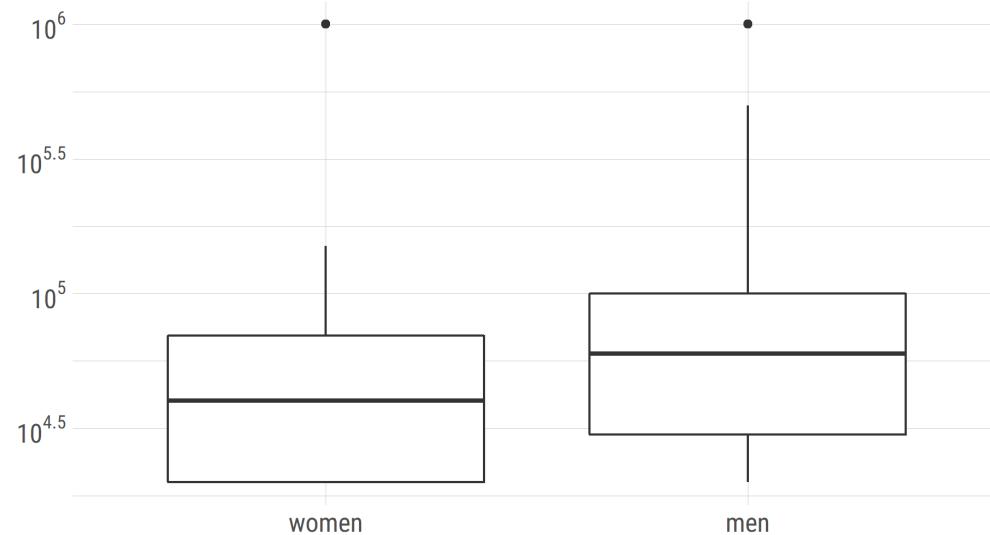
Income is presented as $\log_{10}(\text{Income})$

Box Plot

```
p3 <- okcupid %>%
  slice_sample(n = 5000) %>%
  ggplot(aes(sex, income)) +
  scale_y_log10(breaks = trans_breaks("log10", function(x) 10^x),
                 labels = trans_format("log10", math_format(10^.x)))
  labs(x = "", y = "", title = "OkCupid: Annual Income Dist. by Ge",
       caption = "Income is presented as log10(Income)") +
  scale_x_discrete(labels = c("women", "men")) +
  theme_ipsum(base_family = "Roboto Condensed")

p3 + geom_boxplot()
```

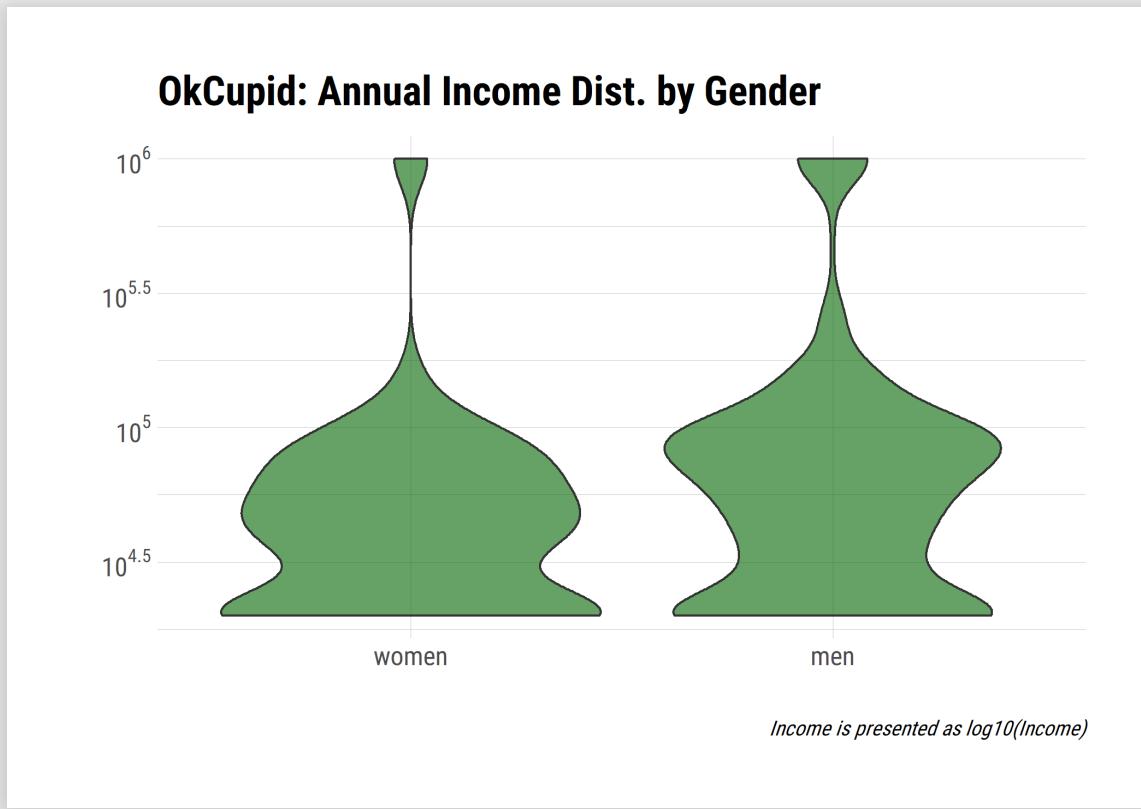
OkCupid: Annual Income Dist. by Gender



Income is presented as $\log_{10}(\text{Income})$

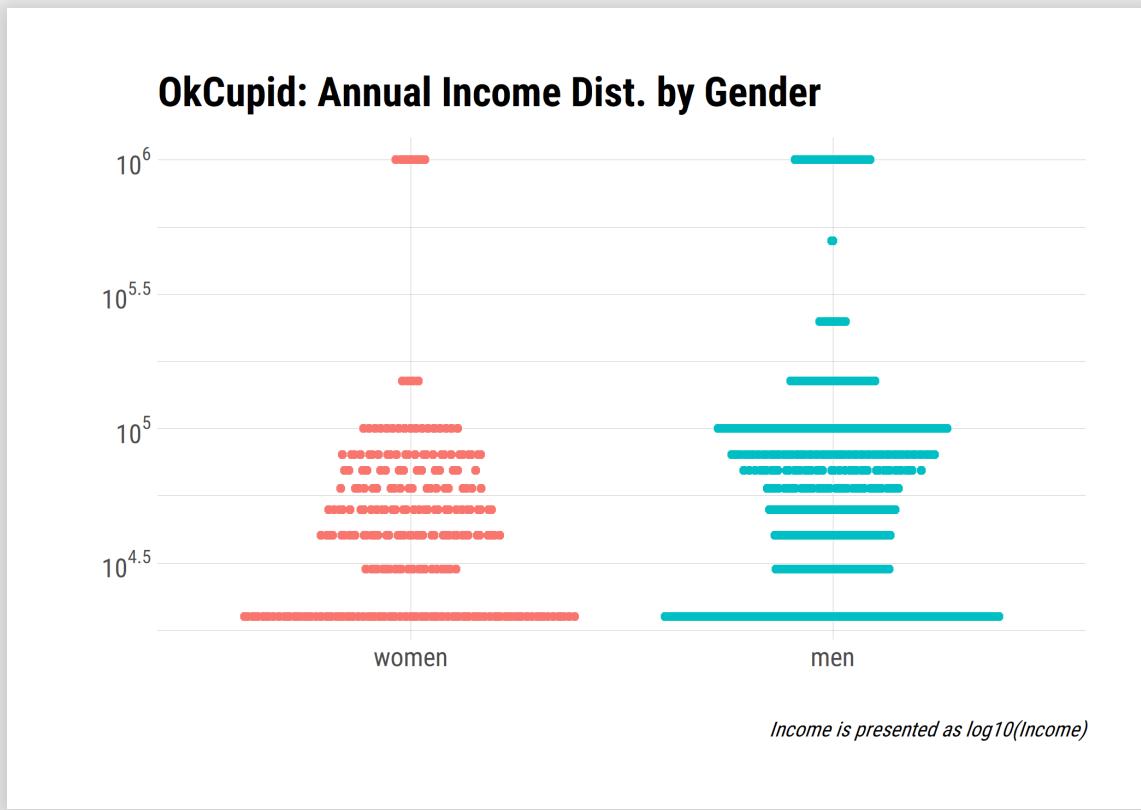
Violin Plot

```
p3 + geom_violin(fill = "darkgreen", alpha = 0.6)
```



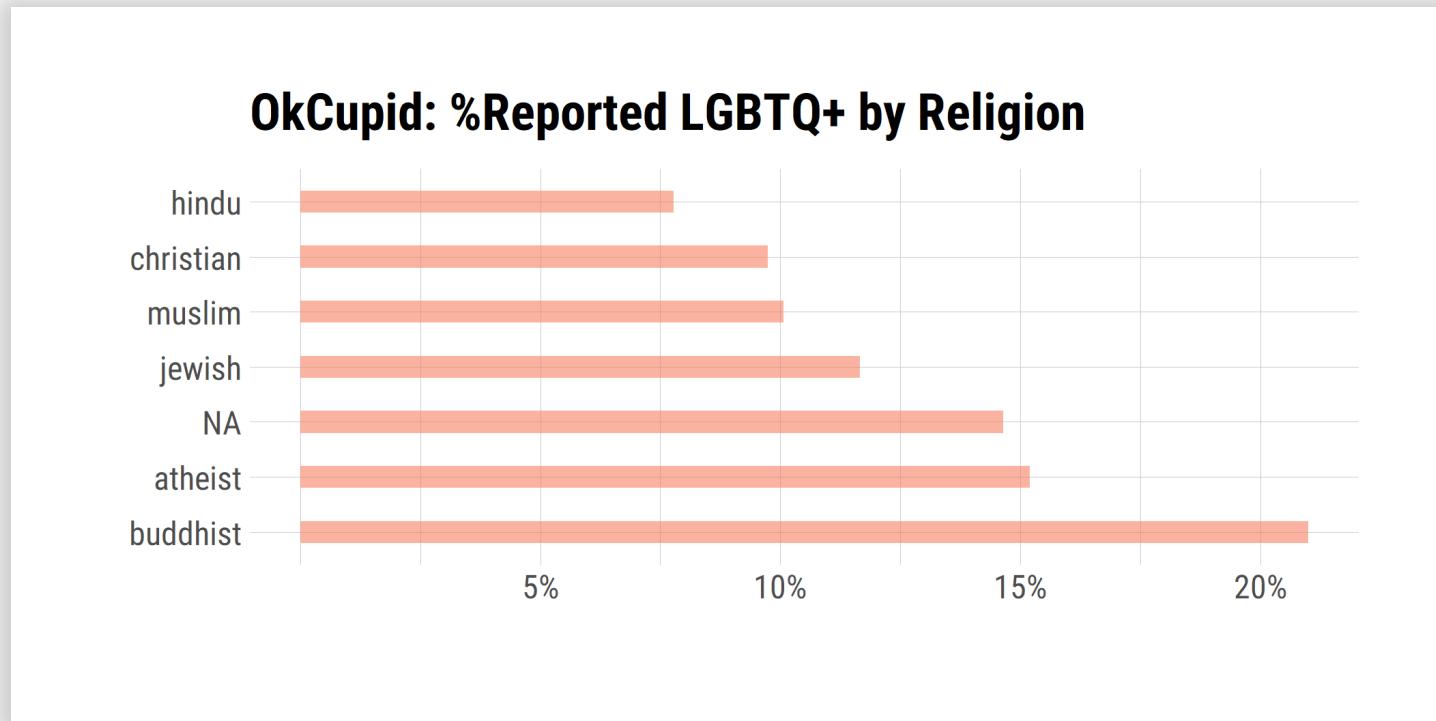
Swarm Plot

```
p3 + ggbeeswarm::geom_quasirandom(aes(color = sex)) + guides(color
```



Bar Plot

```
p + theme_ipsum(base_family = "Roboto Condensed")
```

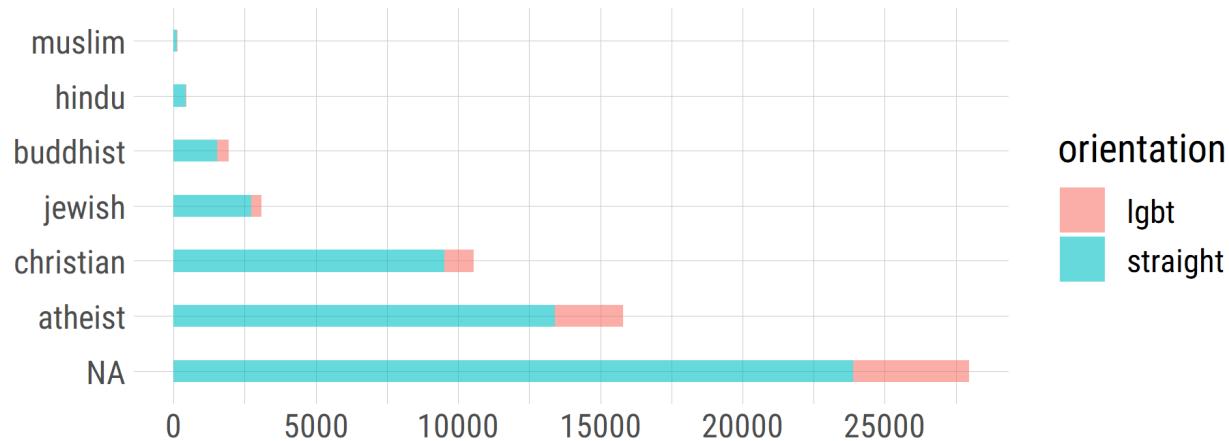


Segmented Bar Plots: Counts

```
p4 <- ggplot(
  okcupid %>%
    mutate(orientation = recode(orientation, "gay" = "lgbt", "bisex" =
      religion = fct_relevel(religion2,
        count(okcupid, religion2, sort =
      aes(religion, fill = orientation)
    ) +
    labs(x = "", y = "",
      title = "OkCupid: Orientation by Religion") +
    scale_y_continuous(breaks = seq(0, 25000, 5000)) +
    coord_flip() +
    theme_ipsum(base_family = "Roboto Condensed") + theme(text = ele
```

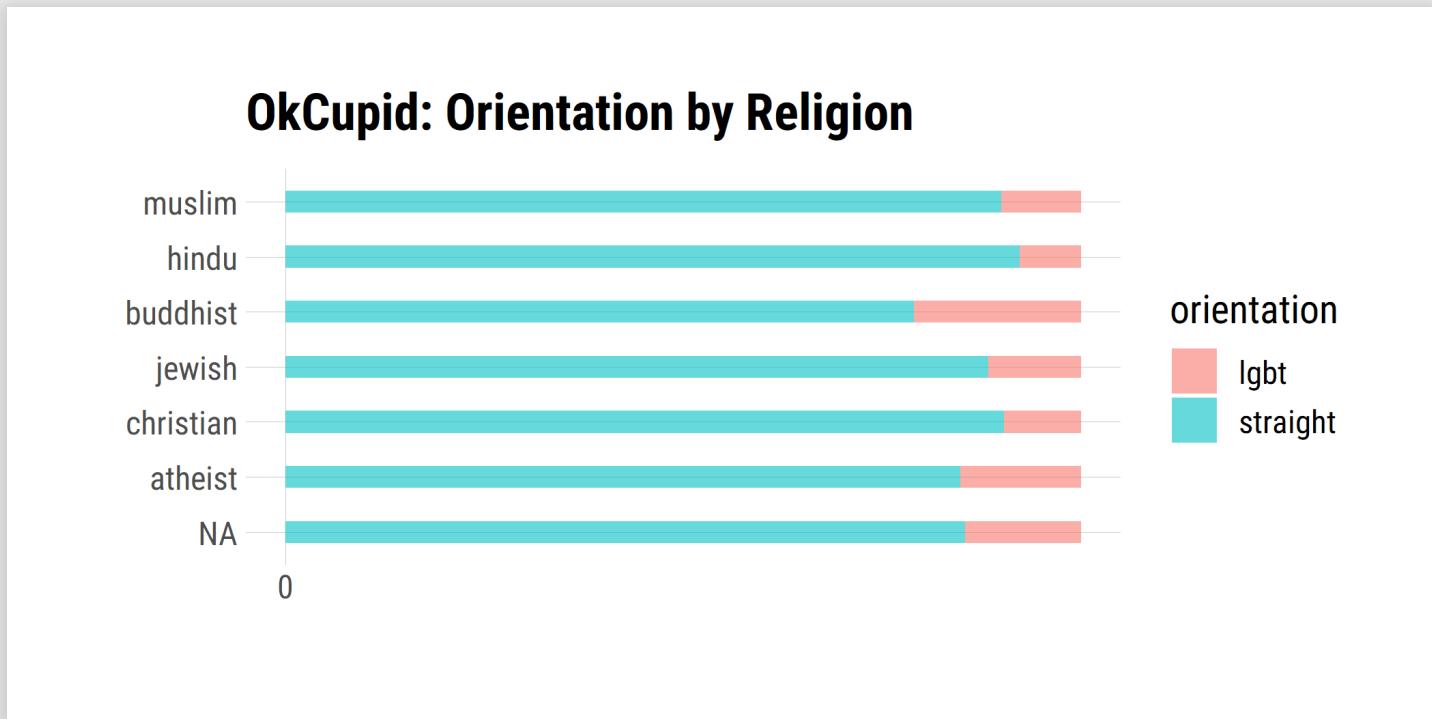
```
p4 + geom_bar(alpha=.6, width=.4)
```

OkCupid: Orientation by Religion



Segmented Bar Plots: Proportions

```
p4 + geom_bar(position = "fill", alpha=.6, width=.4)
```



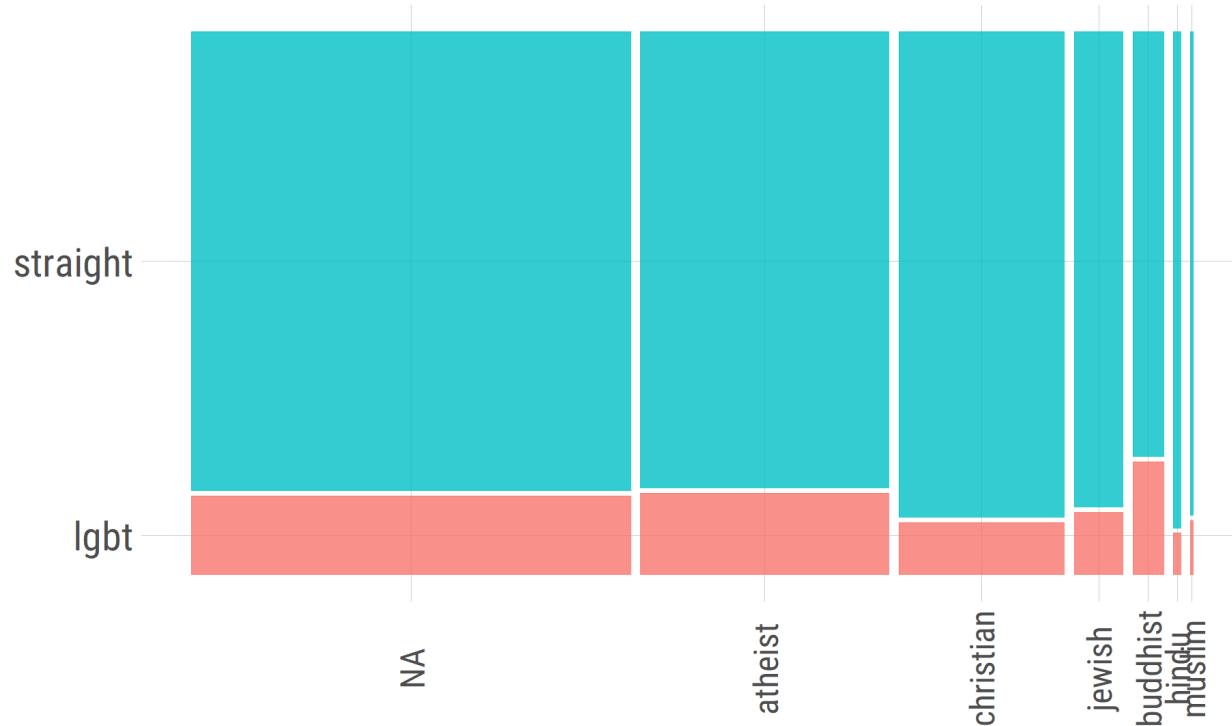
Hoping you can see why this is worse than our original plot.

Mosaic Plots

```
library(ggmosaic)

okcupid %>%
  mutate(orientation = recode(orientation, "gay" = "lgbt", "bisexu
    religion = fct_relevel(religion2, count(okcupid, religior
ggplot() +
  geom_mosaic(aes(x = product(orientation, religion), fill=orienta
  labs(x = "", y = "", title = "OkCupid: Orientation by Religion")
  guides(fill=FALSE) +
  theme_ipsum(base_family = "Roboto Condensed") +
  theme(axis.text.y = element_text(size = 14),
        axis.text.x = element_text(angle = 90, vjust = 0.5, size =
```

OkCupid: Orientation by Religion



Radar Chart

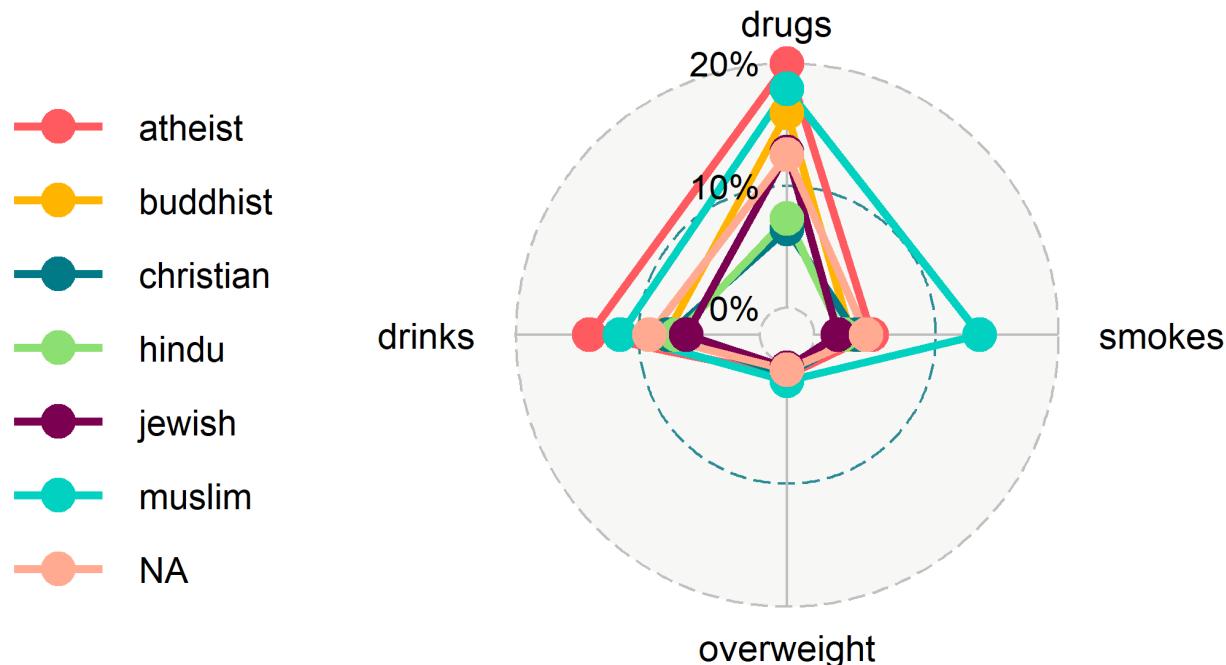
```
okcupid %>%
  group_by(religion2) %>%
  summarise(drugs = mean(drugs %in% c("sometimes", "often"), na.rm = TRUE),
            smokes = mean(smokes == "yes", na.rm = TRUE),
            overweight = mean(body_type == "overweight", na.rm = TRUE),
            drinks = mean(drinks %in% c("often", "very often", "de",
                                         "rately", "never"), na.rm = TRUE))
  ) -> okcupid_radar
okcupid_radar

## # A tibble: 7 x 5
##   religion2  drugs  smokes  overweight  drinks
##   <chr>      <dbl>   <dbl>     <dbl>    <dbl>
## 1 atheist    0.201   0.0473    0.0110   0.141
## 2 buddhist   0.160   0.0302    0.00494  0.0775
## 3 christian  0.0649  0.0352    0.00994  0.0759
## 4 hindu      0.0733  0.0236    0.00725  0.0689
## 5 jewish     0.128   0.0197    0.00467  0.0607
## 6 muslim     0.180   0.136     0.0155   0.115
## 7 NA          0.126   0.0427    0.00634  0.0911
```

```
library(ggradar)

ggradar(okcupid_radar, grid.max = 0.201, grid.mid = 0.10, values.
        plot.title = "OkCupid: Health Indicators by Religion") + t
```

OkCupid: Health Indicators by Religion



Sankey Diagram (Alluvial Plot)

```
library(ggalluvial)

g7 <- c("canada", "france", "germany", "italy", "japan", "united_kingdom"

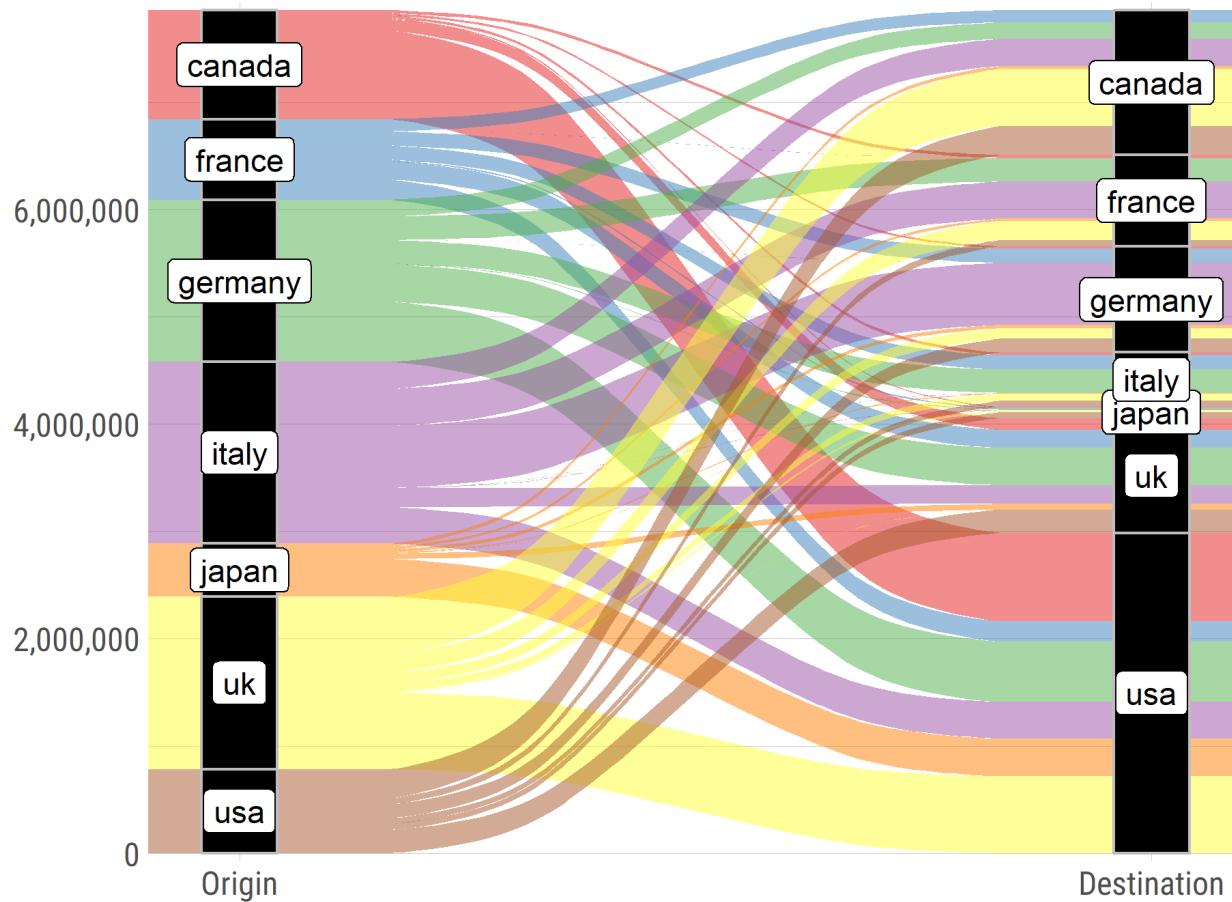
migration_2019 <- read_rds("../data/migration.rds") %>%
  pivot_longer(cols = -c(1:4), names_to = "country_orig", values_to =
  filter(year == 2019) %>%
  filter(across(c("country_orig", "country_dest"), ~.x %in% g7)) %>%
  mutate(across(c("country_orig", "country_dest"), ~case_when(.x =
  group_by(country_orig, country_dest) %>%
  summarise(n_migrants = sum(n_migrants))

migration_2019 %>% head(4)

## # A tibble: 4 x 3
## # Groups:   country_orig [1]
##   country_orig country_dest n_migrants
##   <chr>        <chr>          <dbl>
## 1 canada       canada          0
## 2 canada       france         28339
## 3 canada       germany        18211
## 4 canada       italy          25716
```

```
ggplot(migration_2019,  
       aes(y = n_migrants, axis1 = country_orig, axis2 = country_des)  
     geom_alluvium(aes(fill = country_orig)) +  
     geom_stratum(width = 1/12, fill = "black", color = "grey") +  
     geom_label(stat = "stratum", infer.label = TRUE) +  
     scale_x_discrete(limits = c("Origin", "Destination"), expand = c(0, 0)) +  
     scale_fill_brewer(type = "qual", palette = "Set1") +  
     guides(fill = FALSE) +  
     labs(y = "", title = "G7 Cross-Country Migration 2019") +  
     scale_y_comma() +  
     theme_ipsum(base_family = "Roboto Condensed")
```

G7 Cross-Country Migration 2019



Heatmap

```
gender_age_height_drink <- okcupid %>%
  mutate(age_group = case_when(
    age < 30 ~ "19-29",
    age < 40 ~ "30-39",
    age < 50 ~ "40-49",
    age < 60 ~ "50-59",
    TRUE ~ "60+"),
  height_group = case_when(
    height_cm < 150 ~ "149-",
    height_cm < 160 ~ "150-159",
    height_cm < 170 ~ "160-169",
    height_cm < 180 ~ "170-179",
    height_cm < 190 ~ "180-189",
    TRUE ~ "190+")) %>%
  group_by(sex, age_group, height_group) %>%
  summarise(n = n(),
            drinks = mean(drinks %in% c("often", "very often", "de
```

```
gender_age_height_drink
```

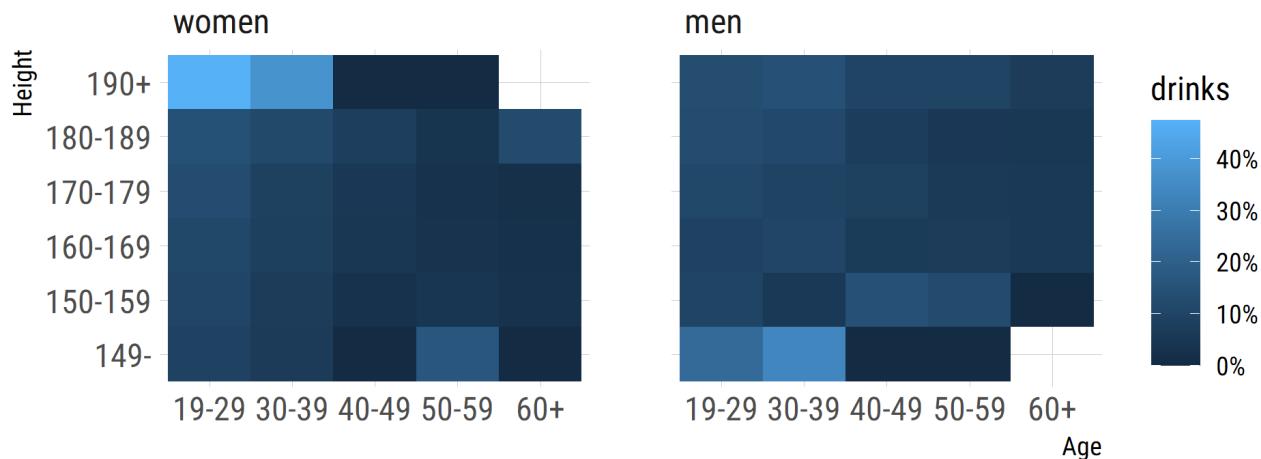
```
## # A tibble: 58 x 5
## # Groups:   sex, age_group [10]
##       sex   age_group height_group      n  drinks
##       <chr>  <chr>        <chr>     <int>    <dbl>
## 1 f       19-29       149-           163  0.0920
## 2 f       19-29       150-159       1963 0.102
## 3 f       19-29       160-169       5620 0.110
## 4 f       19-29       170-179       3113 0.130
## 5 f       19-29       180-189       310  0.145
## 6 f       19-29       190+            19  0.474
## 7 f       30-39       149-           80  0.0625
## 8 f       30-39       150-159       1293 0.0673
## 9 f       30-39       160-169       4109 0.0810
## 10 f      30-39       170-179       2231 0.0852
## # ... with 48 more rows
```

```

heat <- ggplot(gender_age_height_drink, aes(age_group, height_group))
  geom_tile(aes(text = str_c(percent(drinks, 1), " of ", n, " people")))
  facet_wrap(. ~ sex, labeller = sex_labeller) +
  labs(x = "Age", y = "Height", title = "OkCupid: %Reported Drinking by Age and Height")
  scale_fill_gradient(labels = percent_format(accuracy = 1)) +
  theme_ipsum(base_family = "Roboto Condensed")
heat

```

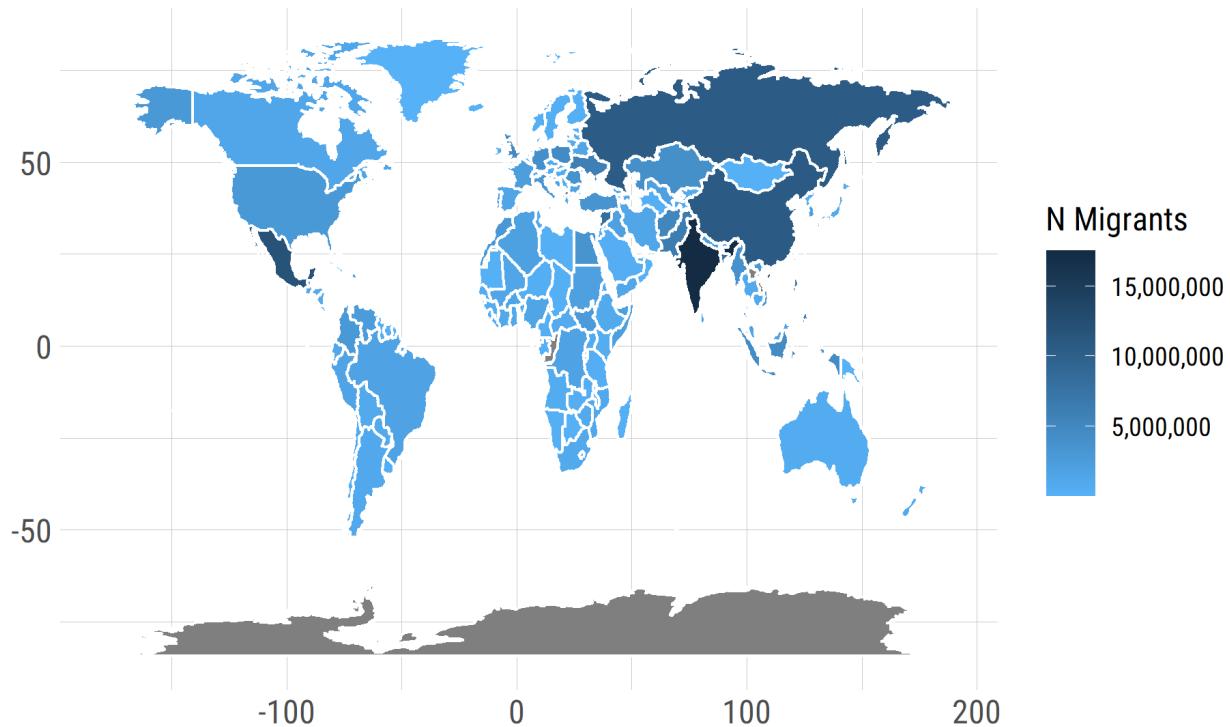
OkCupid: %Reported Drinking by Age and Height



An Actual Heat Map

```
clean_country_name <- function(name) {  
  str_replace_all(str_replace_all(str_to_lower(name), " |-", "_"),  
}  
  
world <- as_tibble(map_data("world"))  
  
world <- world %>%  
  mutate(region = clean_country_name(region))  
  
ggplot(world %>%  
  left_join(migration_total, "region"),  
  aes(x = long, y = lat, group = group)) +  
  geom_polygon(aes(fill = n_migrants), colour = "white") +  
  scale_fill_gradient(labels = comma, name = "N Migrants", high =  
  labs(x = "", y = "", title = "No. of Migrants by Country in 2019")  
  theme_ipsum(base_family = "Roboto Condensed")
```

No. of Migrants by Country in 2019



Amazing Extensions

APPLICATIONS



OF DATA SCIENCE

patchwork

```
library(patchwork)

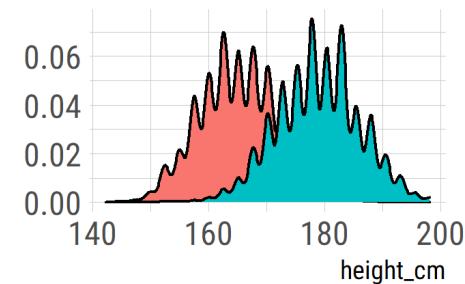
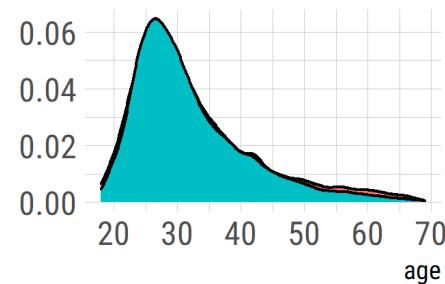
p1 <- ggplot(okcupid %>% filter(between(age, 18, 70))) +
  geom_density(aes(age, fill = sex)) +
  labs(y = "") +
  guides(fill = FALSE) +
  theme_ipsum(base_family = "Roboto Condensed")

p2 <- ggplot(okcupid %>% filter(between(height_cm, 140, 200))) +
  geom_density(aes(height_cm, fill = sex)) +
  labs(y = "") +
  guides(fill = FALSE) +
  theme_ipsum(base_family = "Roboto Condensed")

p3 <- ggplot(okcupid) +
  geom_mosaic(aes(x = product(sex, religion2), fill = sex)) +
  labs(x = "", y = "") +
  theme_ipsum(base_family = "Roboto Condensed") +
  theme(axis.text.x = element_text(angle = 90))
```

Watch this:

```
(p1 | p2) / p3
```

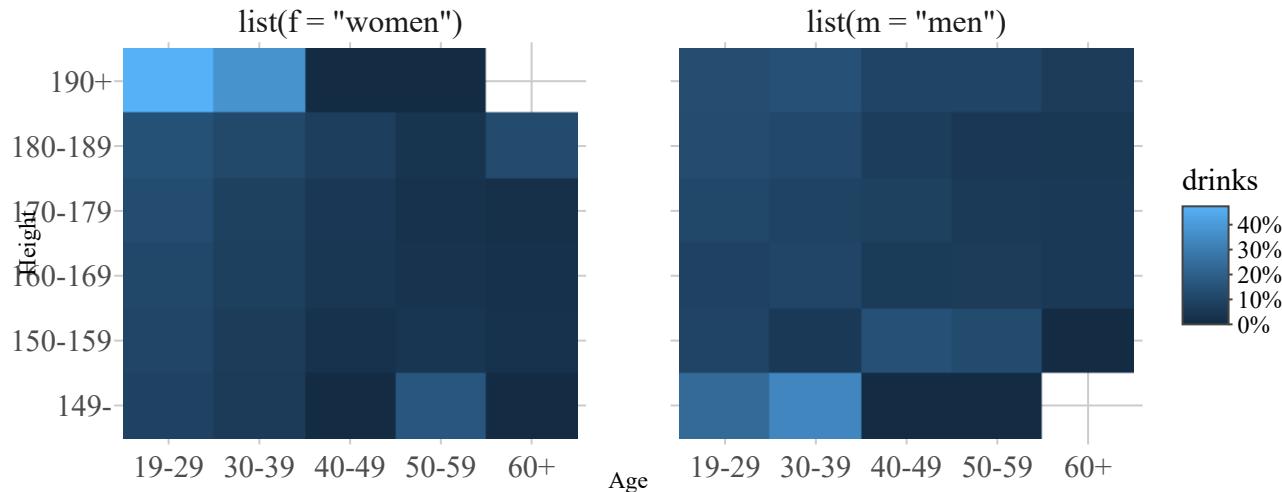


plotly

Remember those heatmaps? They're now interactive!

```
library(plotly)
ggplotly(heat, tooltip = "text")
```

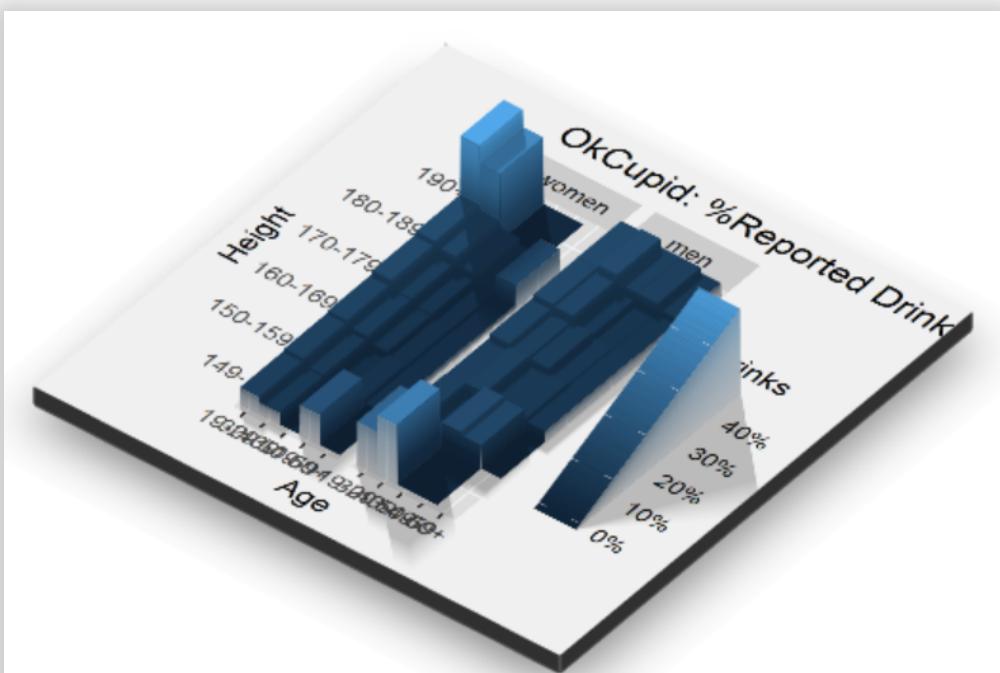
OkCupid: %Reported Drinking by Age and Height



rayshader

Or rotatable 3D...

```
library(rayshader)  
plot_gg(heat)
```



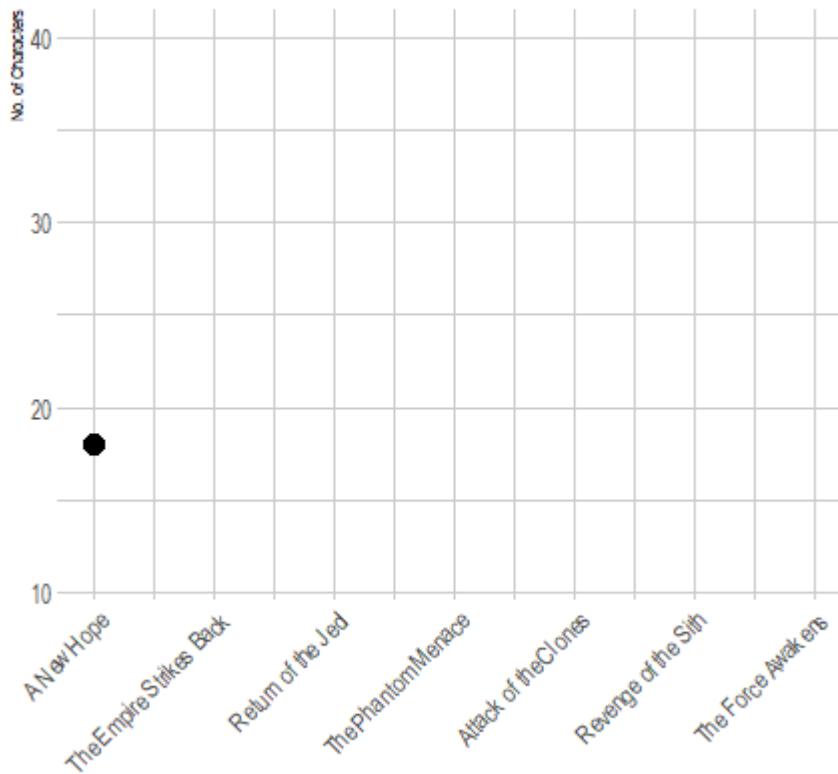
ganimate

Take a plot and make it better:

```
library(gganimate)

characters %>%
  group_by(film_id) %>%
  summarise(n_char = n_distinct(character_id)) %>%
  ggplot(aes(film_id, n_char)) +
  geom_line(lwd = 2) +
  geom_point(aes(group = seq_along(film_id)), size = 5) +
  labs(x = "", y = "No. of Characters",
       title = "No. of Unique Characters: {films$title[frame %% 1]}",
       scale_x_continuous(breaks = 1:7, labels = films$title) +
  theme_ipsum(base_family = "Roboto Condensed") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1)) +
  transition_reveal(along = film_id) +
  ease_aes("cubic-in-out")
```

No. of Unique Characters: A New Hope

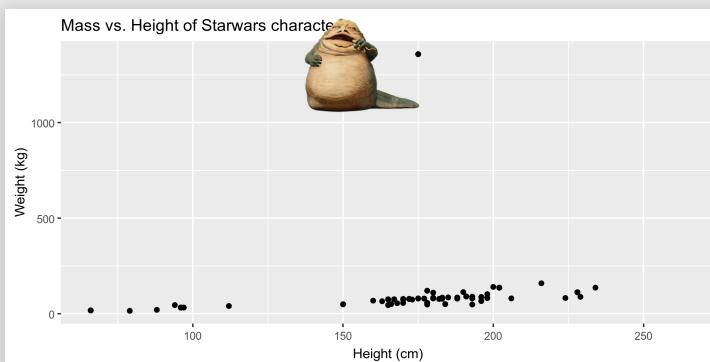


magick

That Jabba Plot was NOT made by Photoshop...

```
library(magick)

jabba <- image_read("images/jabba.png")
fig <- image_graph(width = 2400, height = 1200, res = 300)
ggplot(data = characters, mapping = aes(x = height, y = mass)) +
  geom_point(size = 1.5) +
  labs(title = "Mass vs. Height of Starwars characters",
       x = "Height (cm)", y = "Weight (kg)")
fig %>% image_composite(jabba, offset = "+1000+30")
```



tidygraph + ggraph

```
library(tidygraph)

characters_2_plus <- characters %>%
  add_count(name) %>%
  filter(n > 1) %>%
  select(name, film_id) %>%
  mutate(id = group_indices(., name)) # for some reason 1:n ID

nodes <- characters_2_plus %>%
  count(id, name)

edges <- characters_2_plus %>%
  inner_join(characters_2_plus, "film_id") %>%
  rename(from = id.x,
         to = id.y) %>%
  count(from, to) %>%
  rename(weight = n) %>%
  filter(from != to)

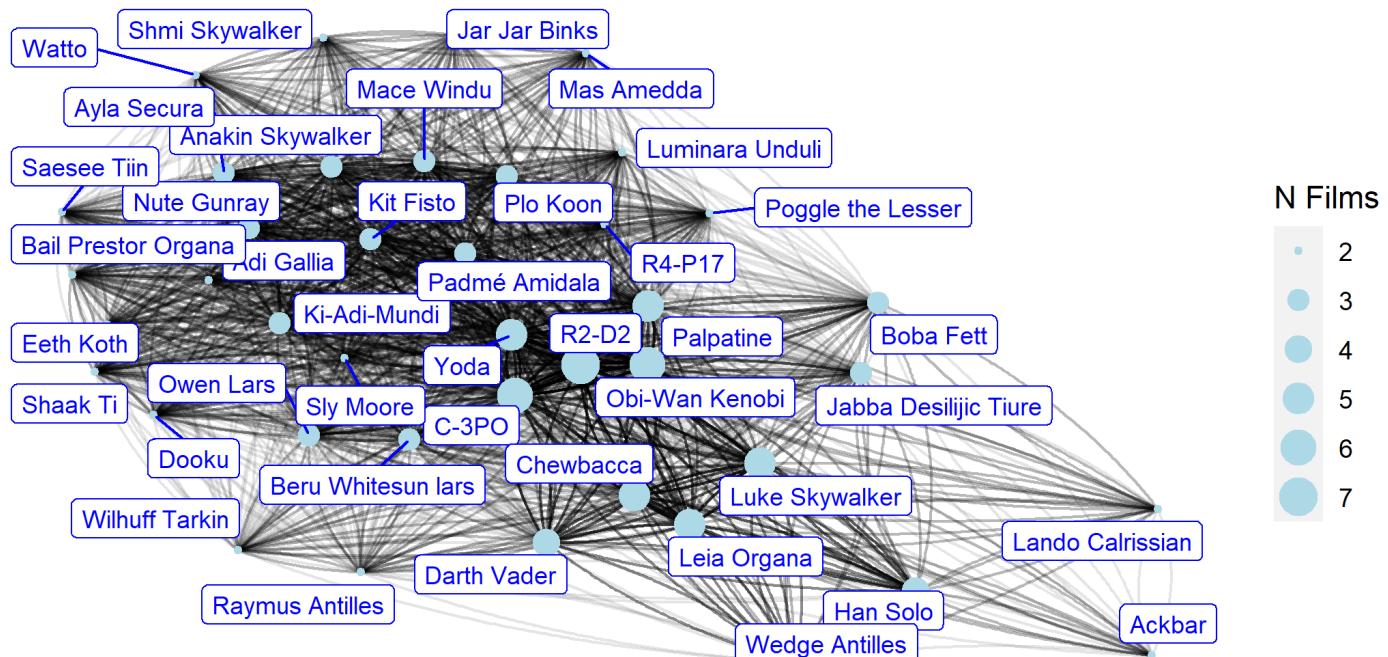
graph <- tbl_graph(nodes = nodes, edges = edges, directed = FALSE)
```

```

library(ggraph)
ggraph(graph, layout = "fr") +
  geom_edge_fan(aes(alpha = weight), show.legend = FALSE) +
  geom_node_point(aes(size = n), color = "lightblue") +
  geom_node_label(aes(label = name), color = "blue", repel = TRUE,
  theme(panel.background = element_rect(fill="white", colour = "white"),
  labs(size = "N Films", title = "Star Wars: Characters Network",

```

Star Wars: Characters Network



An Edge between two characters means they appeared in the same film

Edge transparency means how many films

ggwithimages

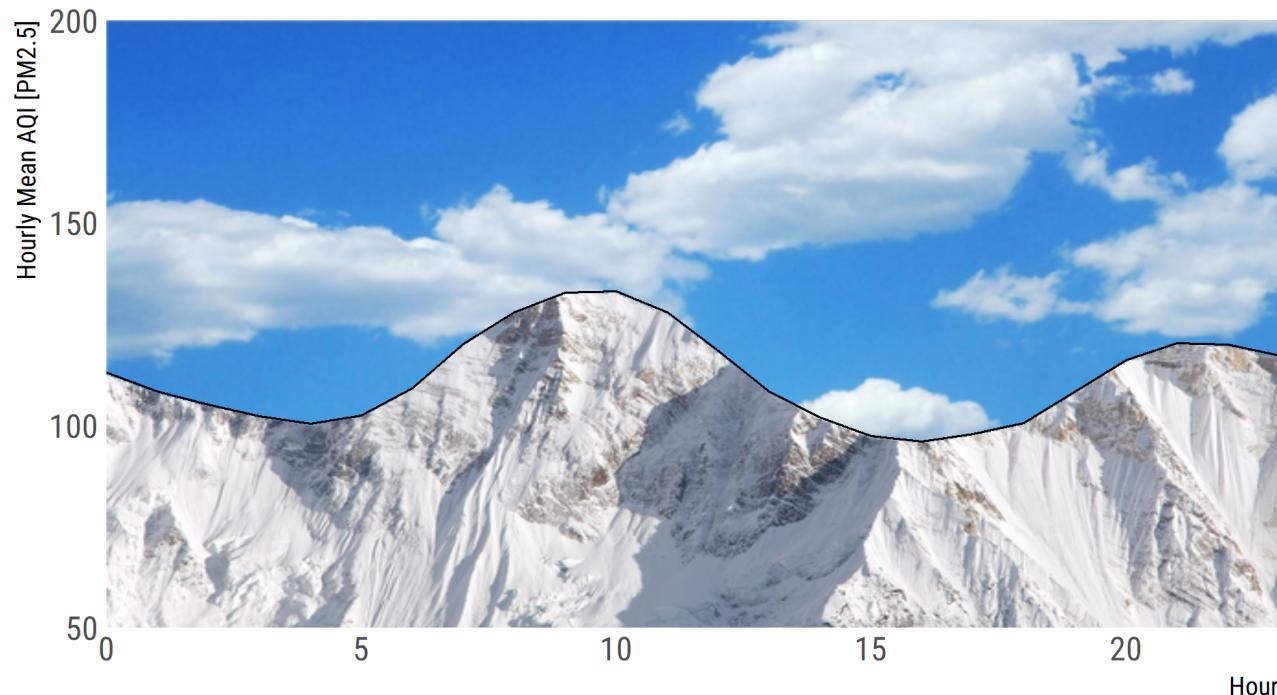
```
library(ggwithimages)

annapurna <- png:::readPNG(system.file("extdata", "annapurna.png",
sky <- png:::readPNG(system.file("extdata", "sky.png", package =
kathmandu_hourly_aqi <- readr:::read_csv(system.file("extdata", "ka

ggplot(kathmandu_hourly_aqi, aes(hour, aqi)) +
  geom_line_with_image(annapurna, sky) +
  labs(title = "Air Quality Index in the Thamel, Kathmandu, Nepal",
       subtitle = "Measured in PM2.5 by the US Embassy in Kathmandu",
       y = "Hourly Mean AQI [PM2.5]",
       x = "Hour") +
  ylim(c(50, 200)) +
  theme_ipsum(base_family = "Roboto Condensed")
```

Air Quality Index in the Thamel, Kathmandu, Nepal

Measured in PM2.5 by the US Embassy in Kathmandu



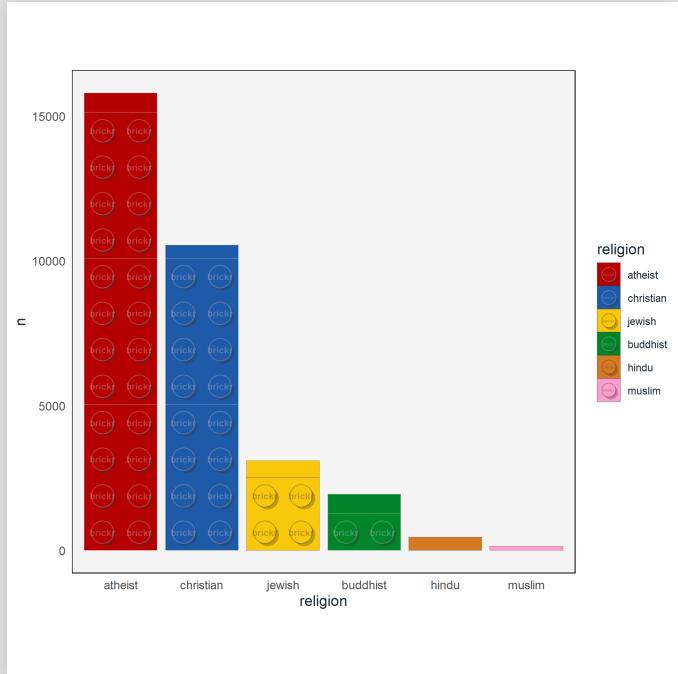
brickr...

```
library(brickr)

okcupid %>%
  count(religion2) %>%
  filter(religion2 != "NA") %>%
  mutate(religion =
    fct_relevel(religion2, c("atheist", "christian", "jewis
ggplot(aes(religion, n)) +
  geom_brick_col(aes(fill = religion)) +
  scale_fill_brick() +
  coord_brick() +
  theme_brick()
```

(Though this was available only on v0.2.0 (H))

And... it can get a lot more fun:



Get the legos you need with
`build_instructions()` 😊